

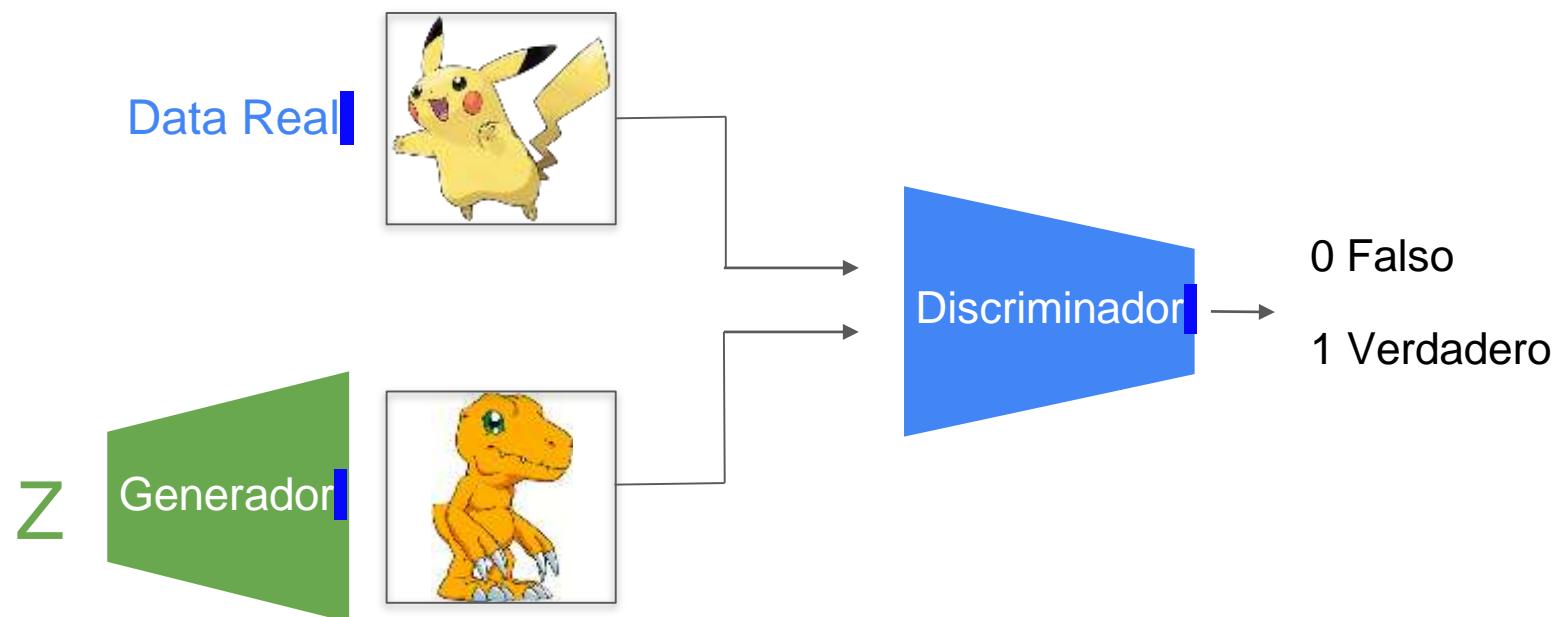
Deep Learning en Imágenes con GANs y Modelos de Difusión - Parte II

Prof. Peter Montalvo

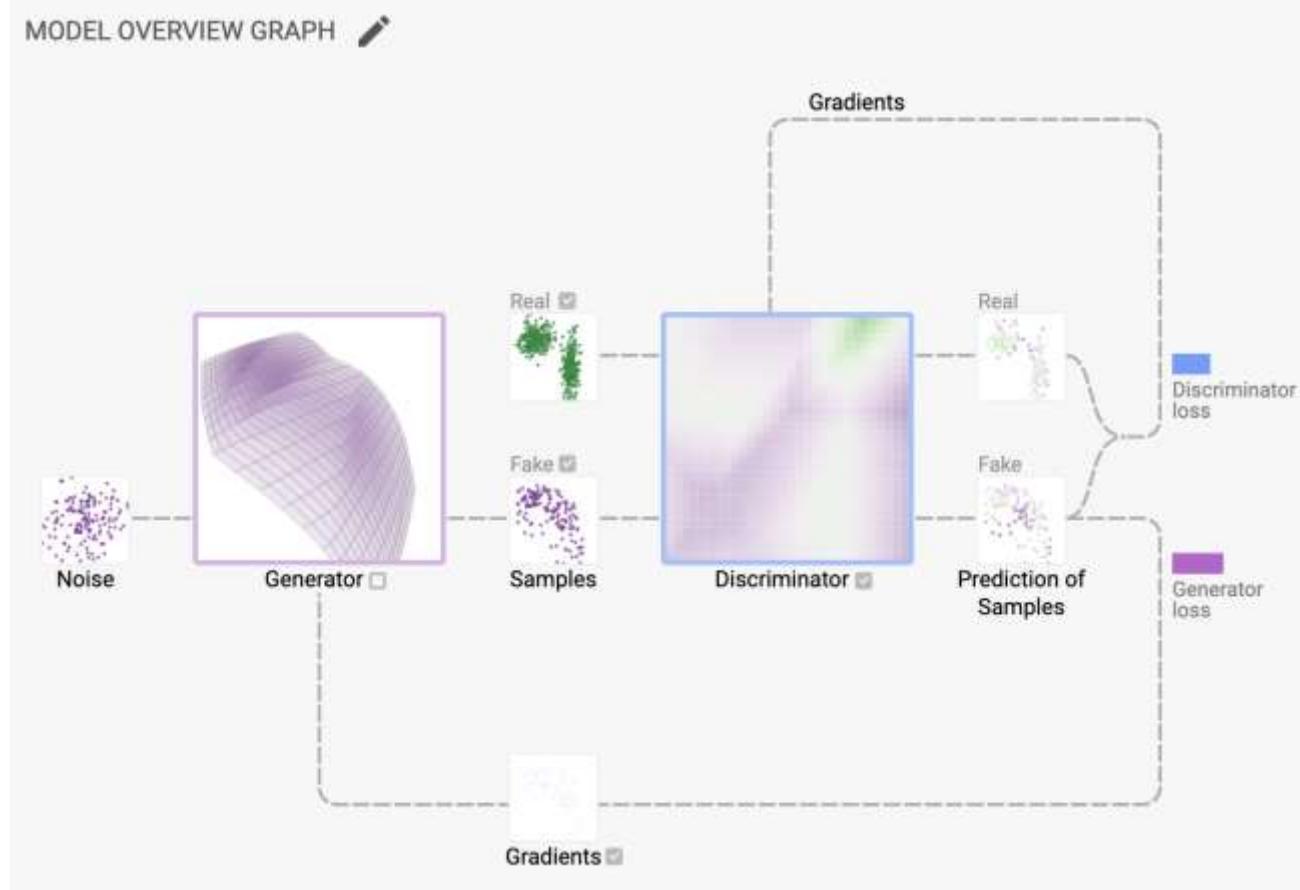
Agenda

- Introducción
- GANs
- Función de Pérdida
- Aplicaciones
- Wasserstein GAN

GAN - Generative Adversarial Nets

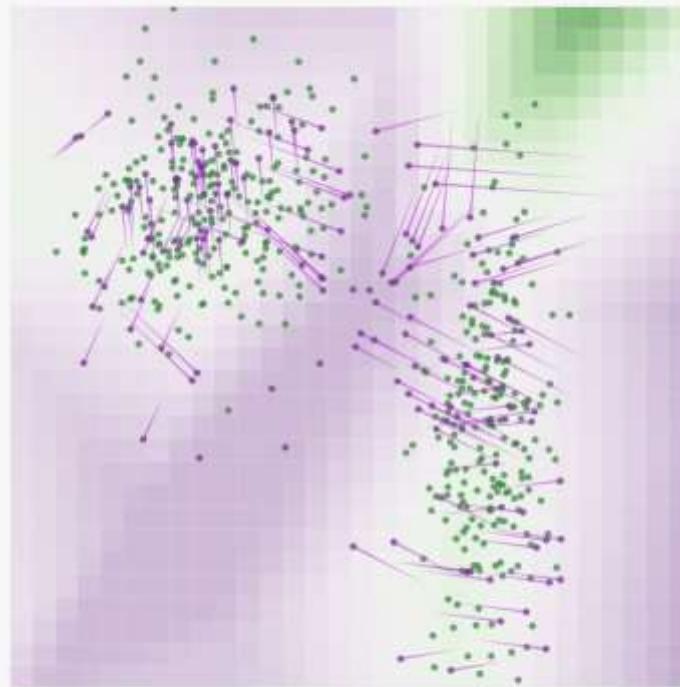


Introducción

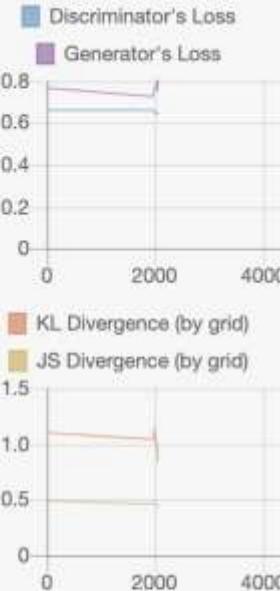


Introducción

LAYERED DISTRIBUTIONS



METRICS



Background colors of grid cells represent **discriminator**'s classifications.

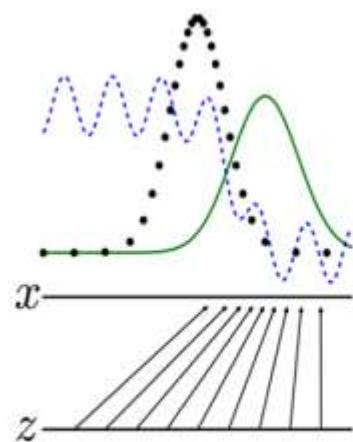
Samples in **green regions** are likely to be real; those in **purple regions** likely fake.

Manifold represents **generator**'s transformation results from noise space.
Opacity encodes density: darker purple means more samples in smaller area.

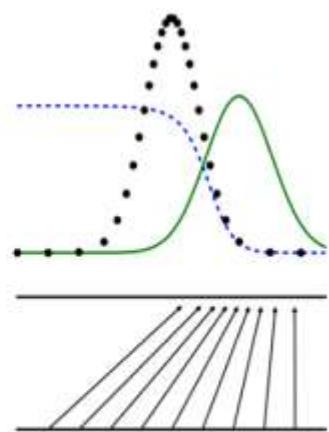
Pink lines from fake samples represent **gradients** for generator.

✓ This sample needs to move upper right to decrease generator's loss.

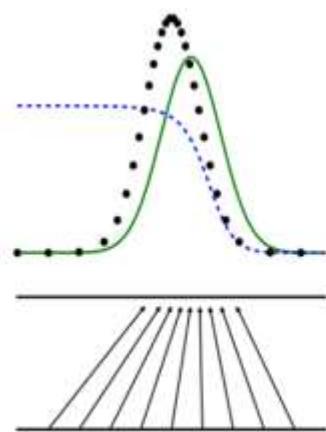
Función de pérdida



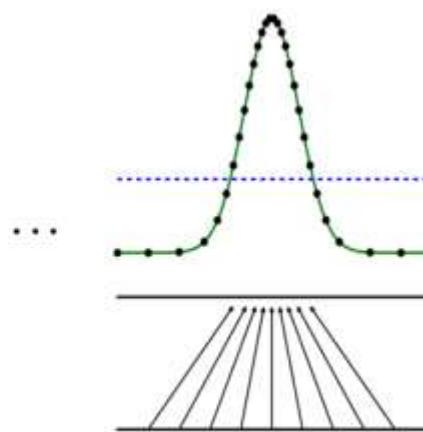
(a)



(b)

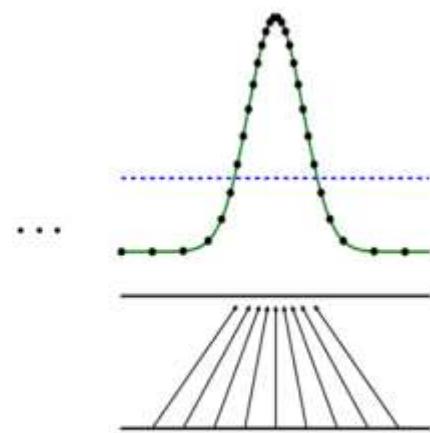
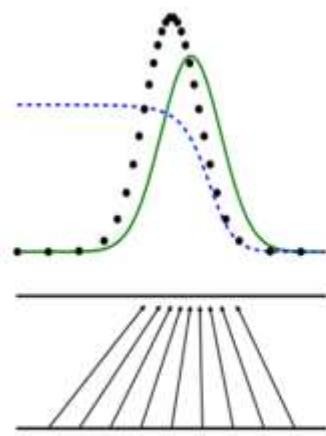
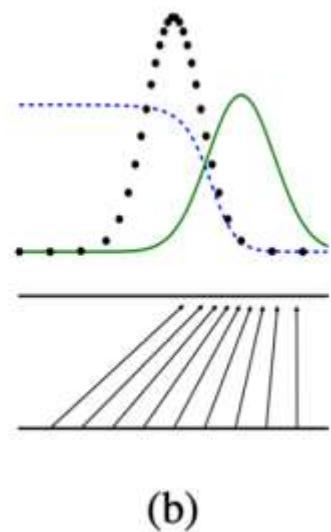
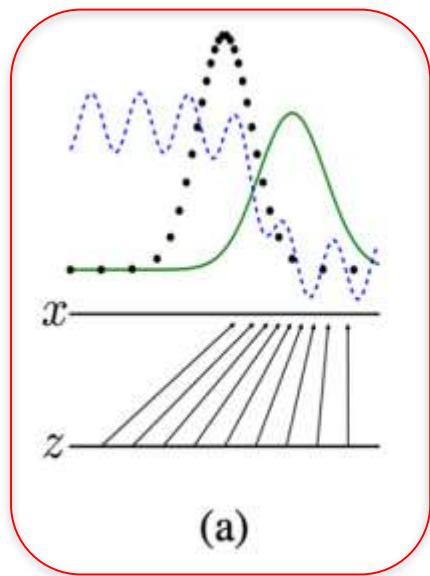


(c)



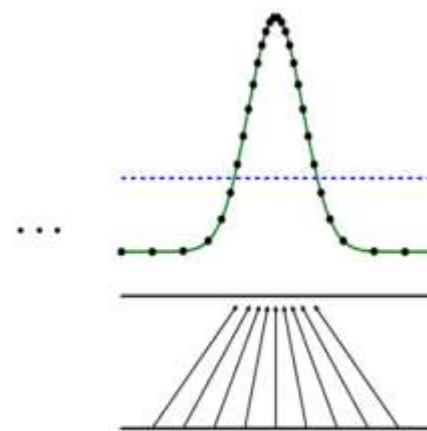
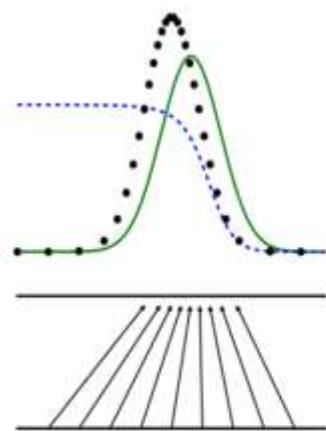
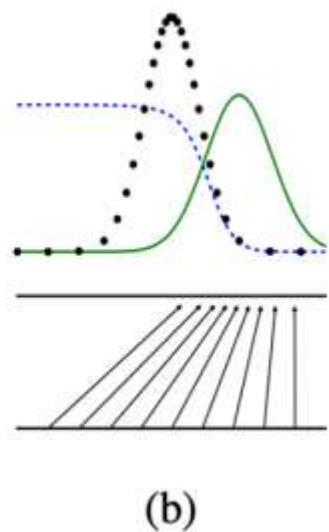
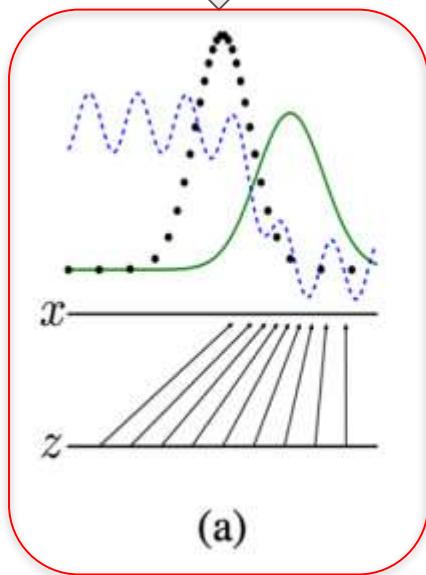
(d)

Función de pérdida

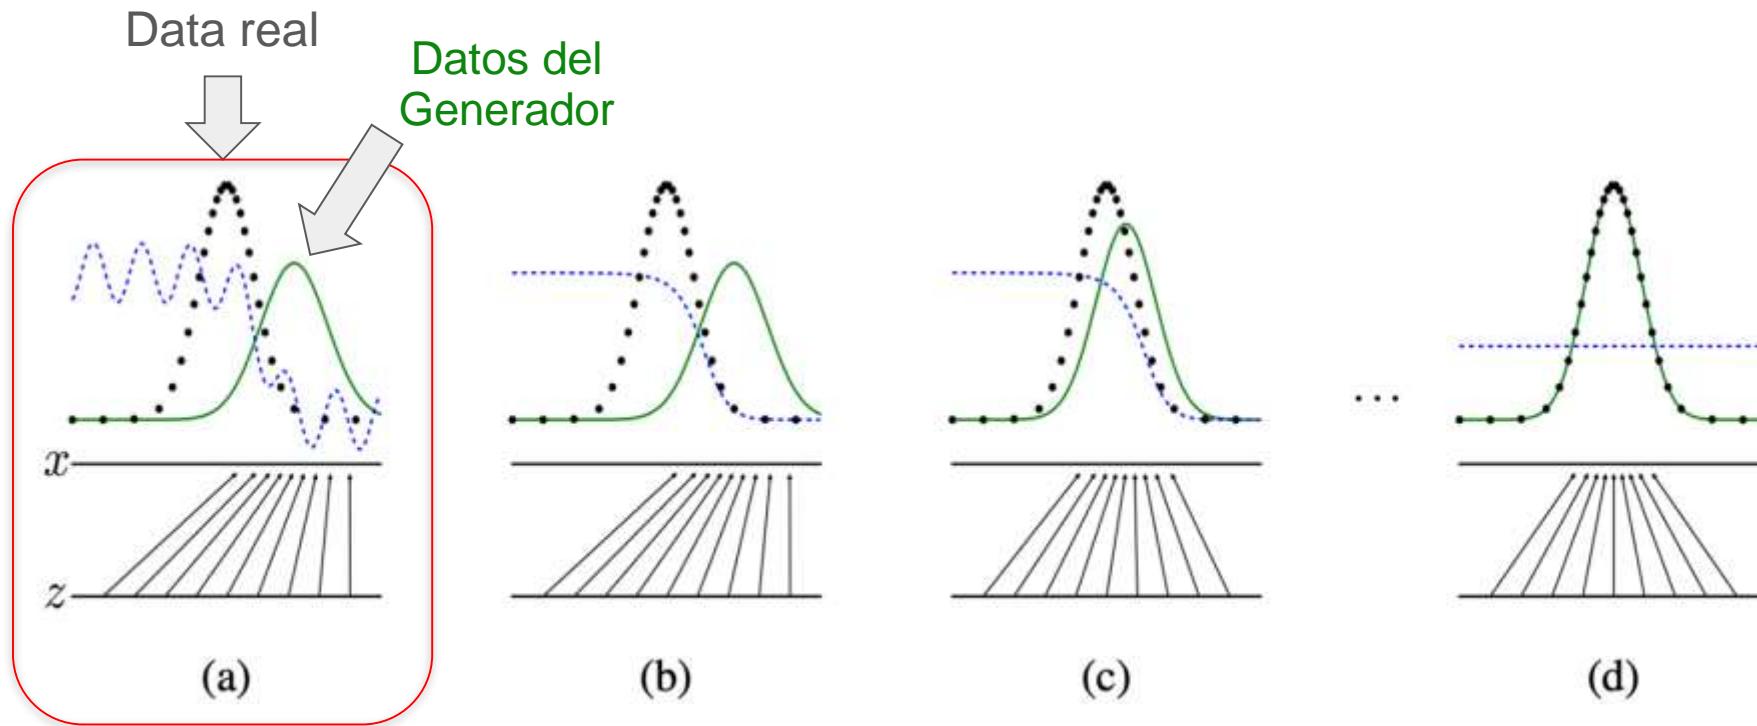


Función de pérdida

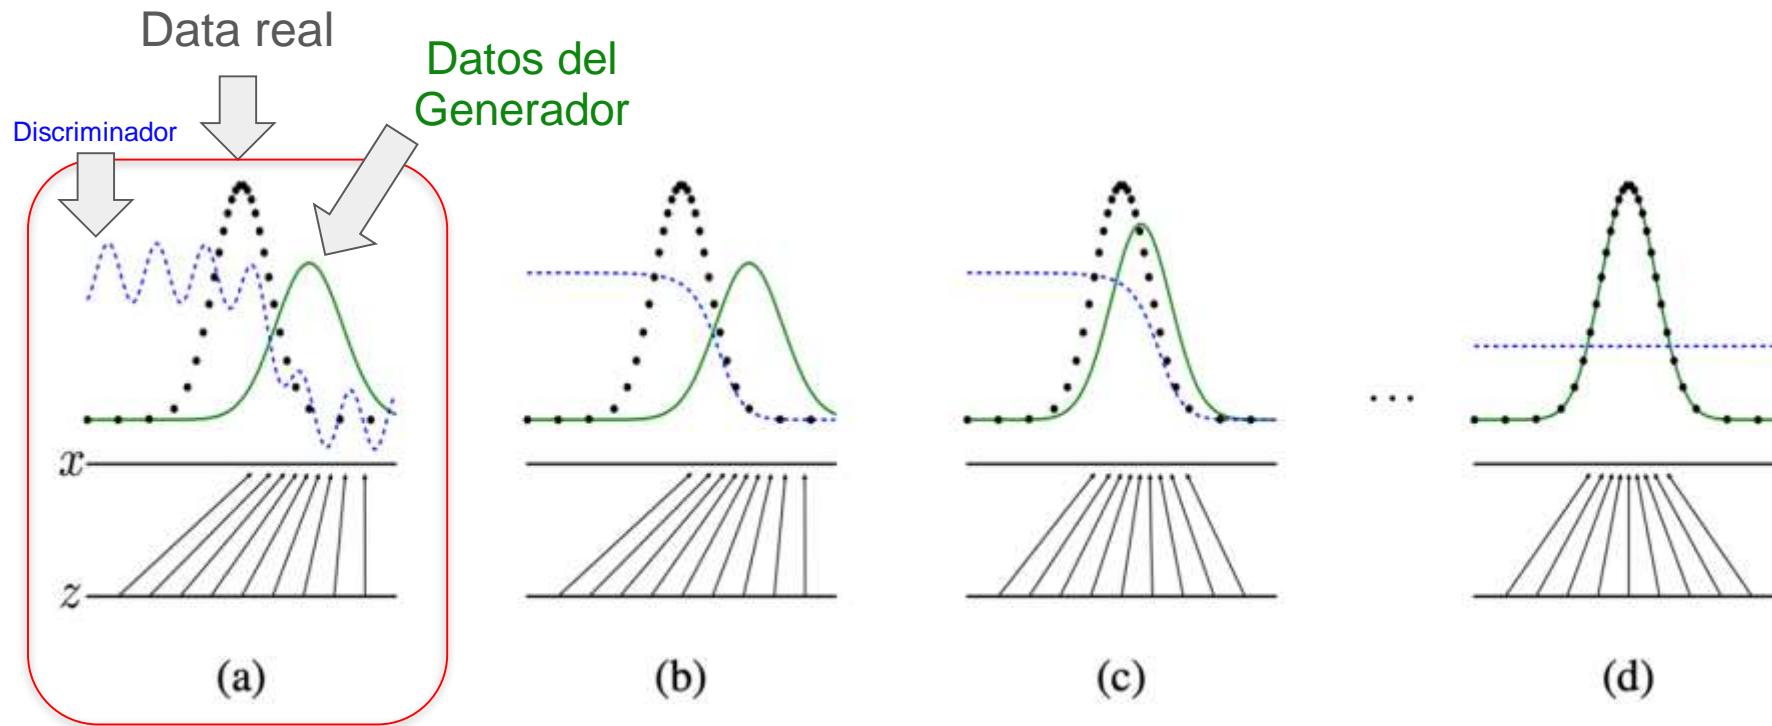
Data real



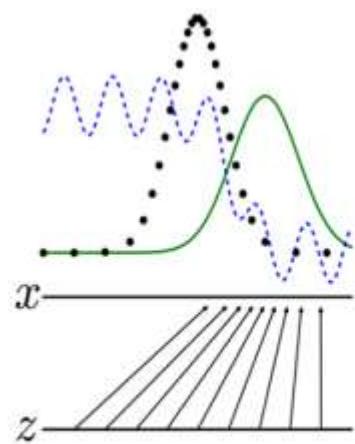
Función de pérdida



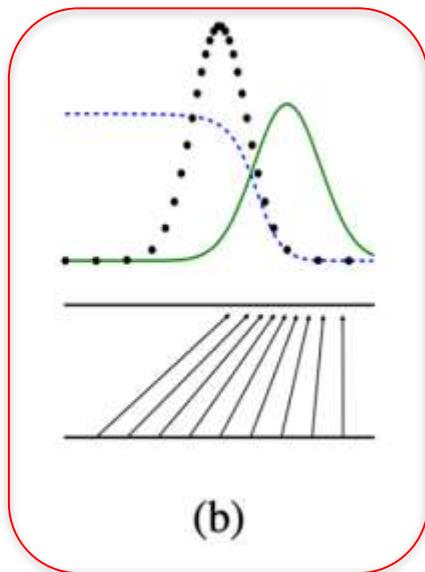
Función de pérdida



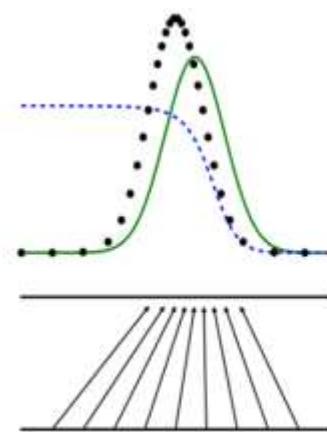
Función de pérdida



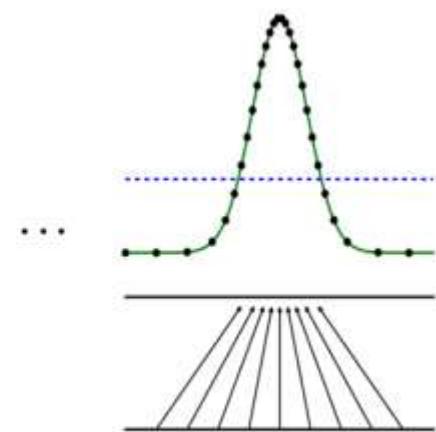
(a)



(b)

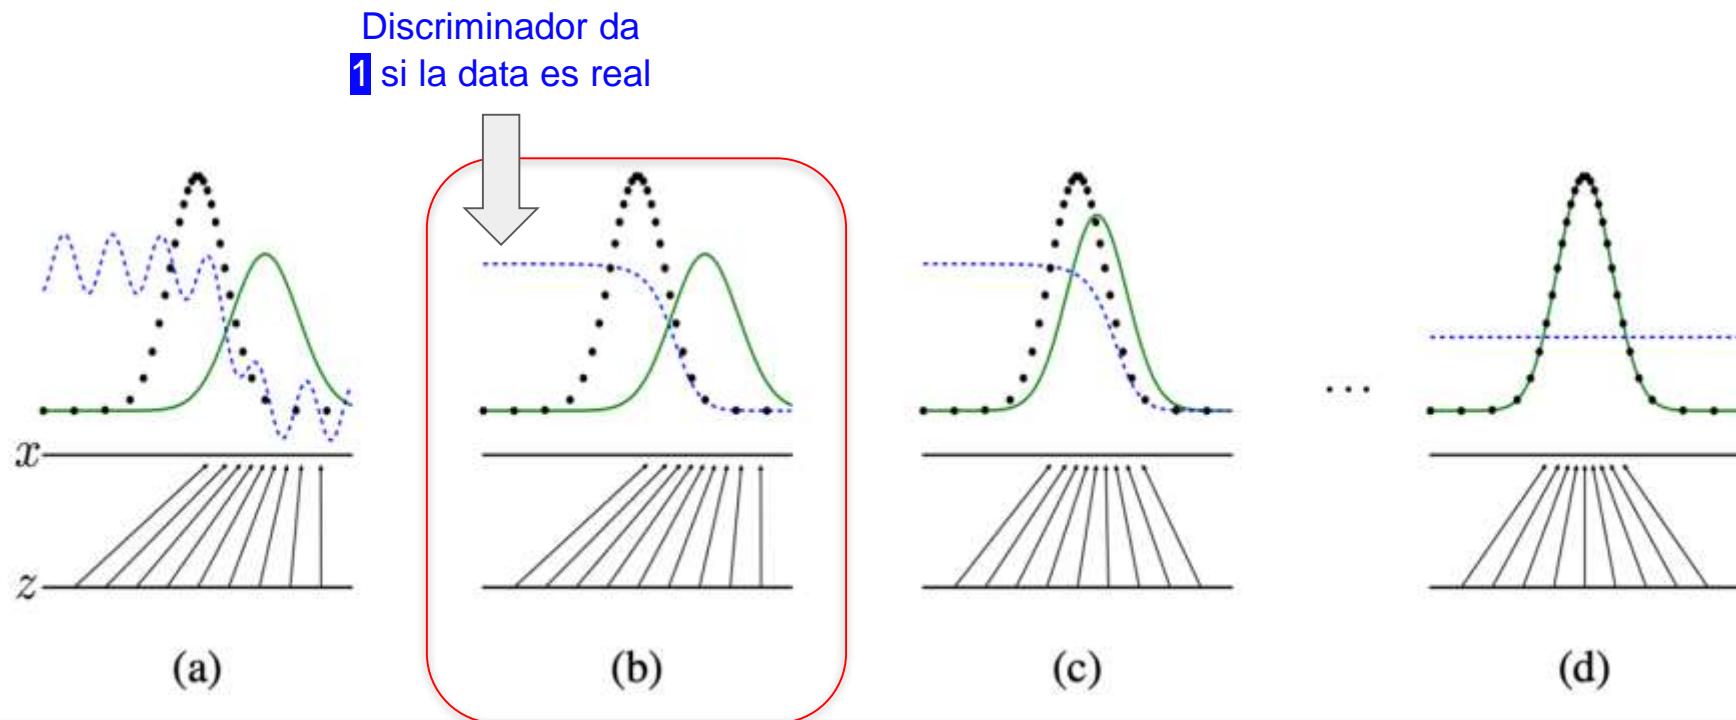


(c)

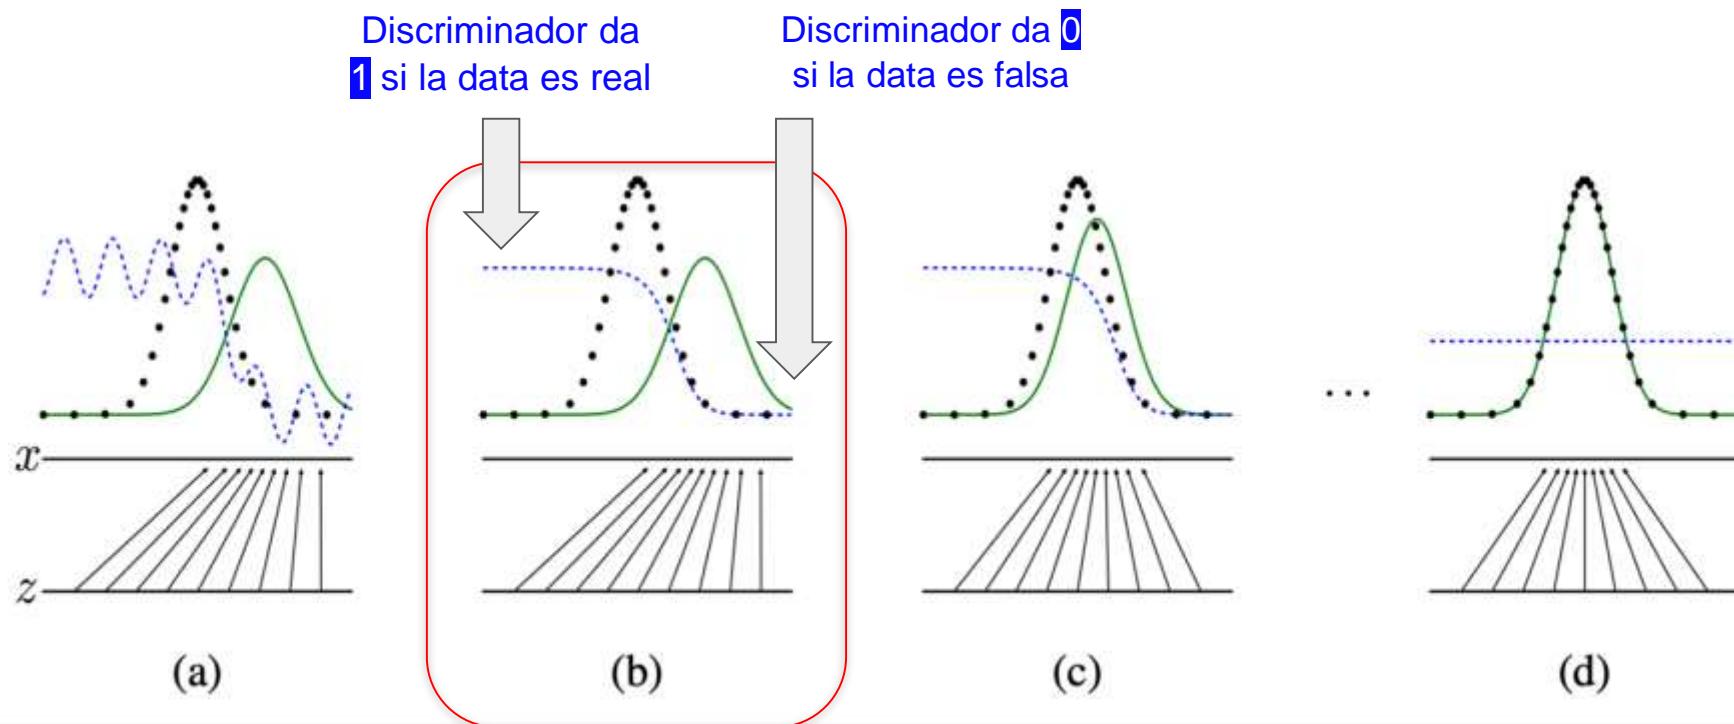


(d)

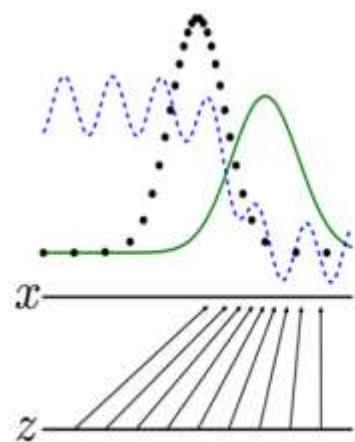
Función de pérdida



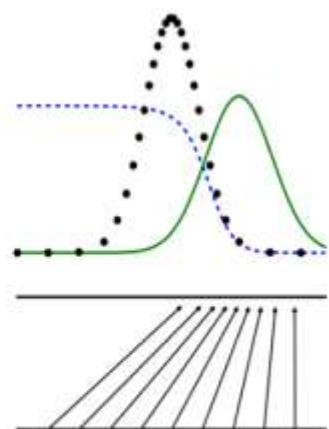
Función de pérdida



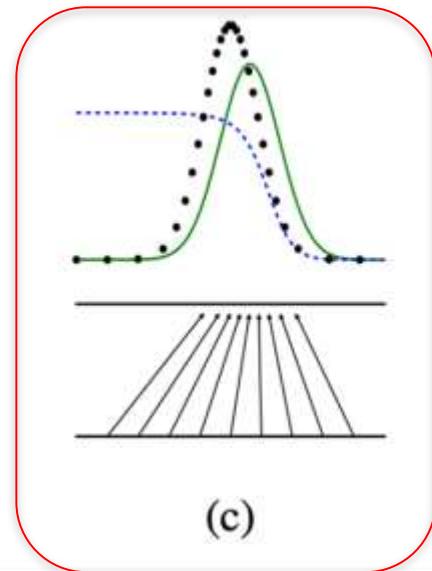
Función de pérdida



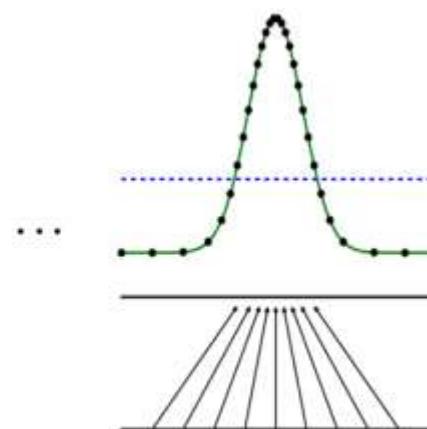
(a)



(b)

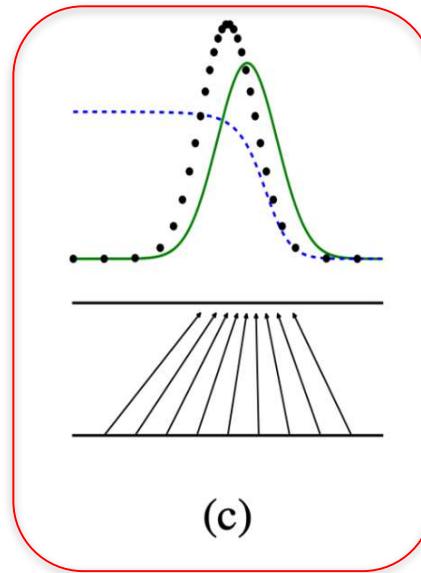


(c)



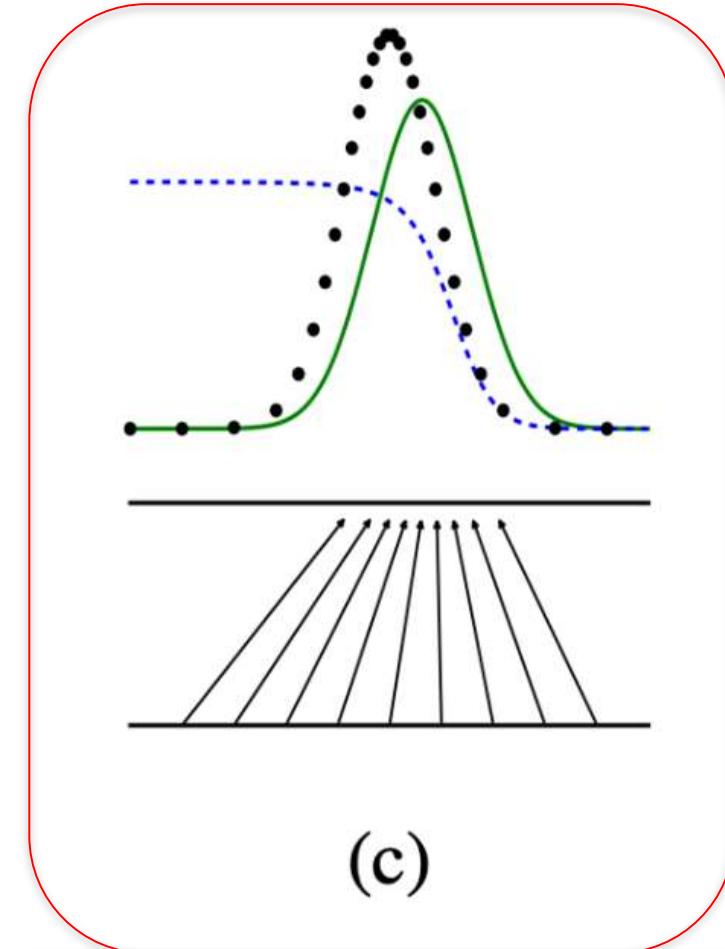
(d)

Función de pérdida



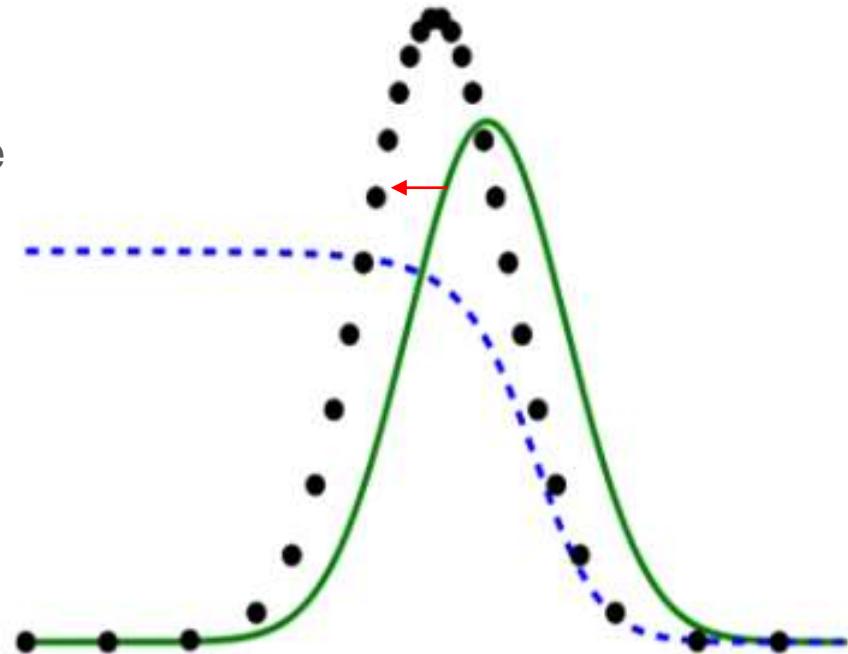
Función de pérdida

La distribución de datos falsos se comienza a acercar a la distribución de datos original



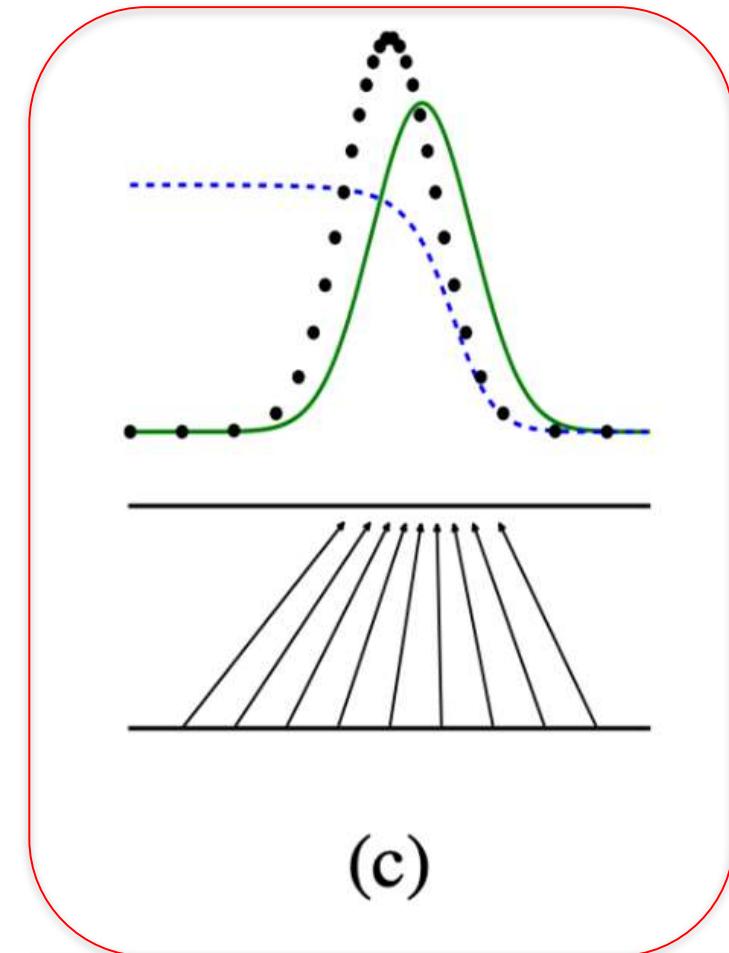
Función de pérdida

La distribución de datos falsos se comienza a acercar a la distribución de datos original



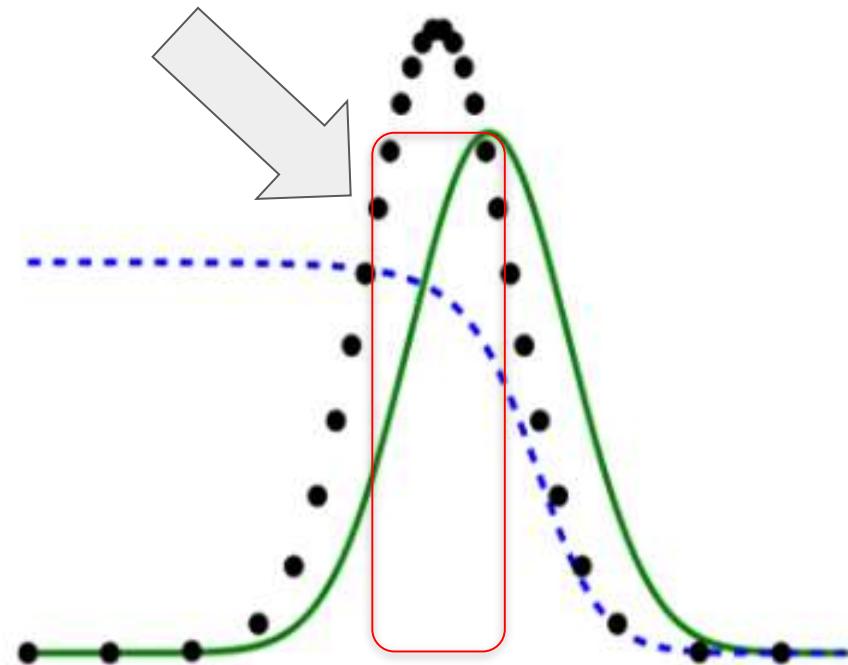
Función de pérdida

El discriminador comienza a clasificar como 1 a la distribución de data falsa.

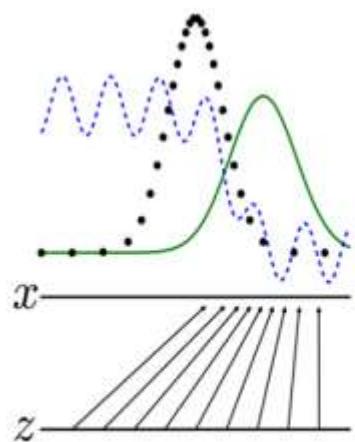


Función de pérdida

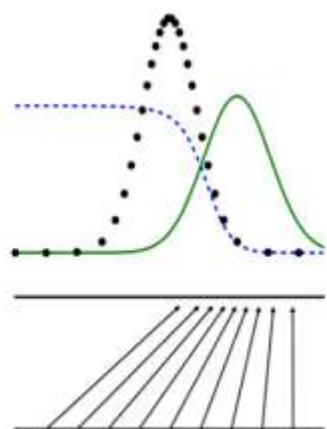
El discriminador comienza a clasificar como 1 a la distribución de data falsa.



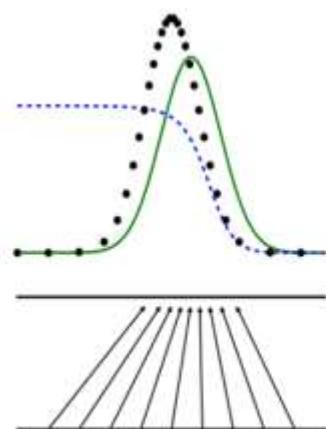
Función de pérdida



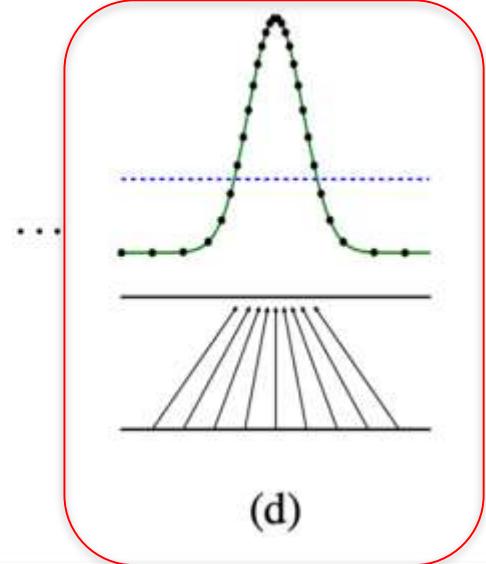
(a)



(b)

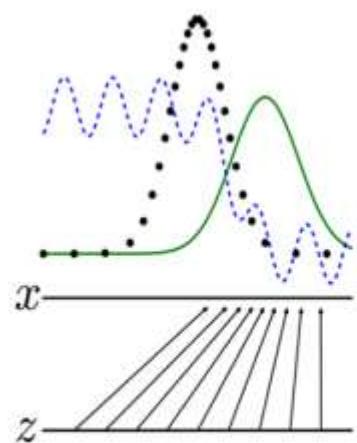


(c)

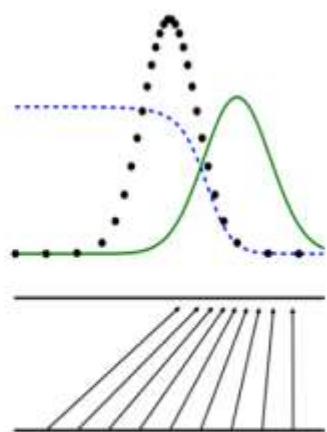


(d)

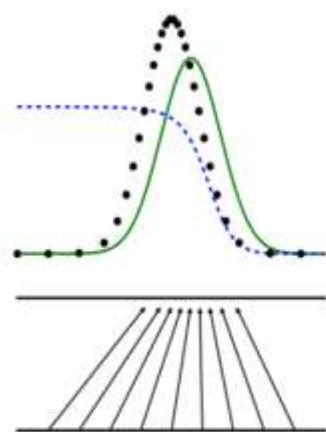
Función de pérdida



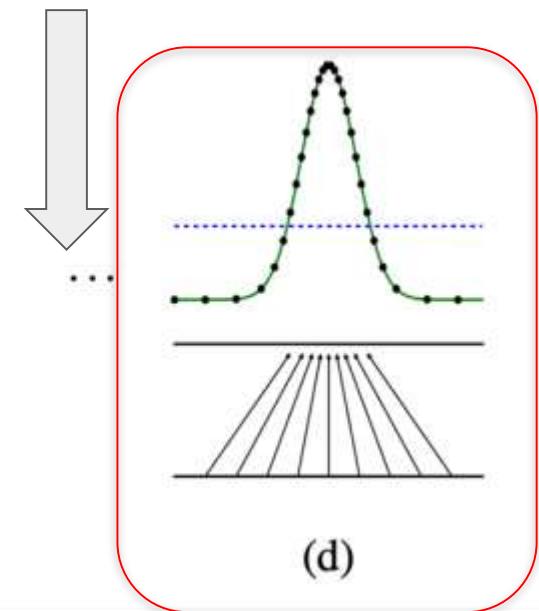
(a)



(b)

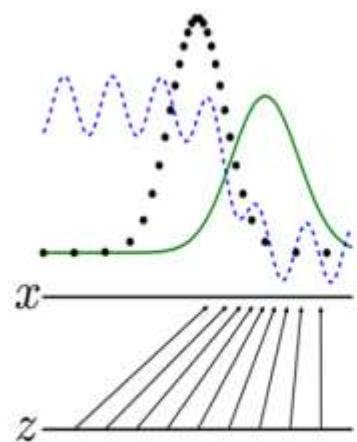


(c)

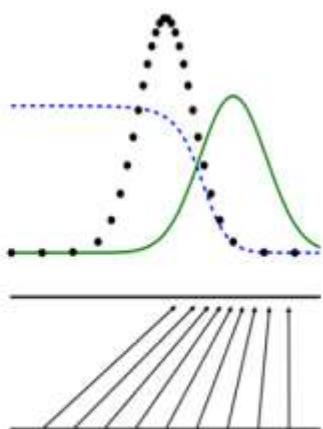


(d)

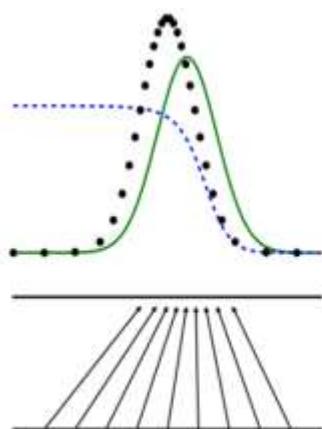
Función de pérdida



(a)

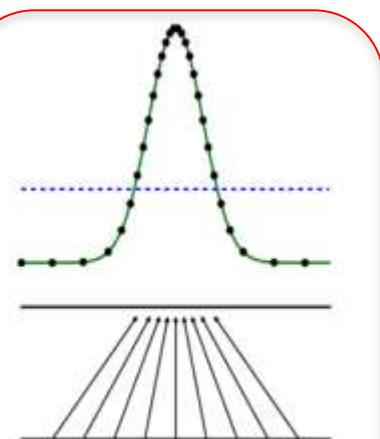
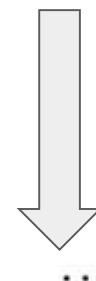


(b)



(c)

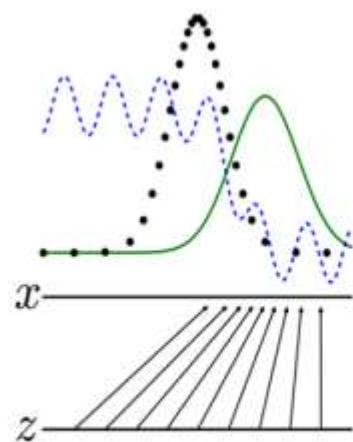
Finalmente, luego de algunas iteraciones...



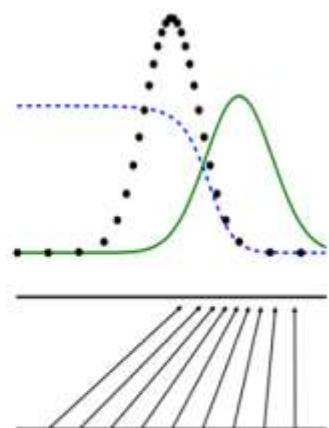
(d)

Función de pérdida

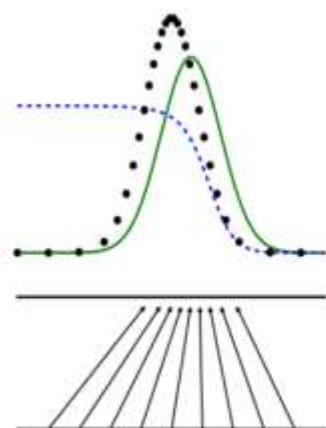
El Generador converge generando data similar a la distribución de data real.



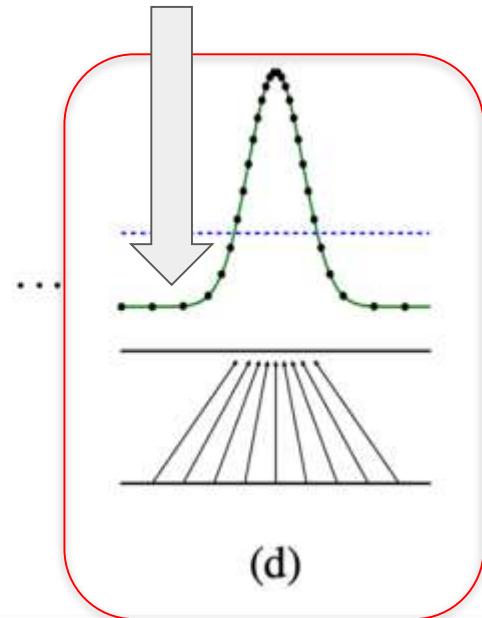
(a)



(b)



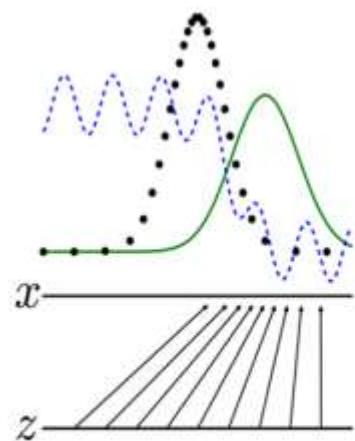
(c)



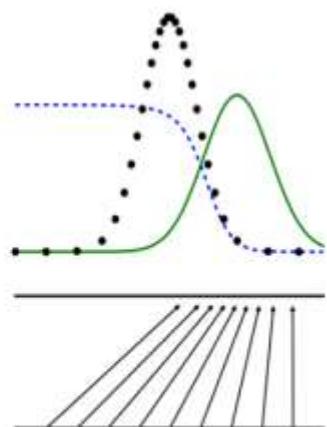
(d)

Función de pérdida

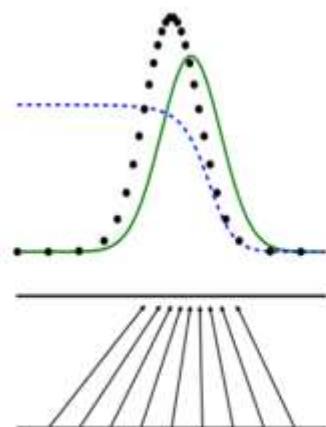
El discriminador converge en 0.5



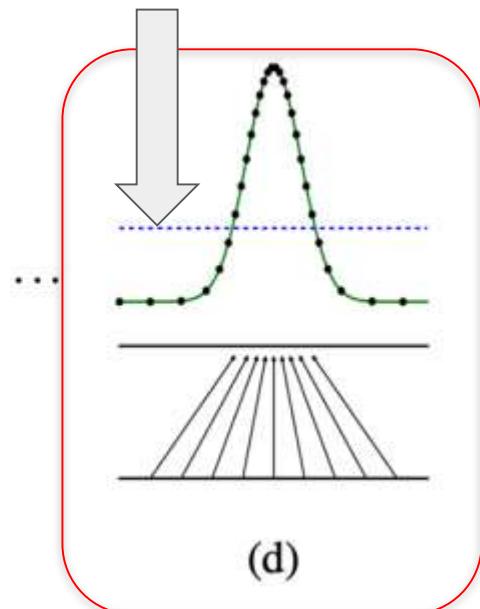
(a)



(b)



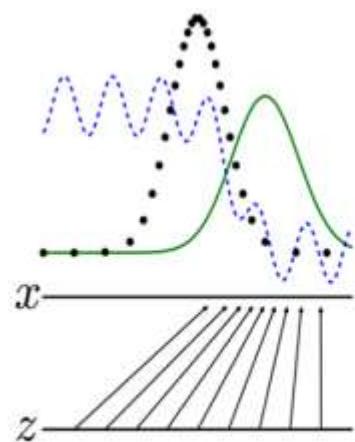
(c)



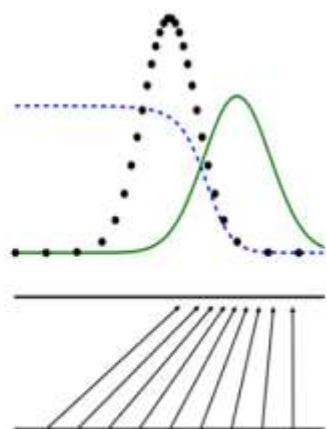
(d)

Función de pérdida

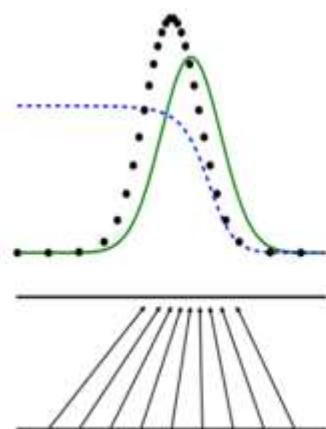
El generador converge en 0.5 (No puede diferenciar entre real y falso)



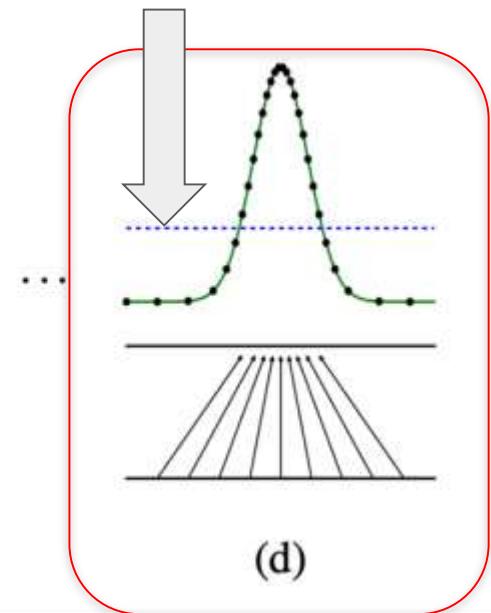
(a)



(b)

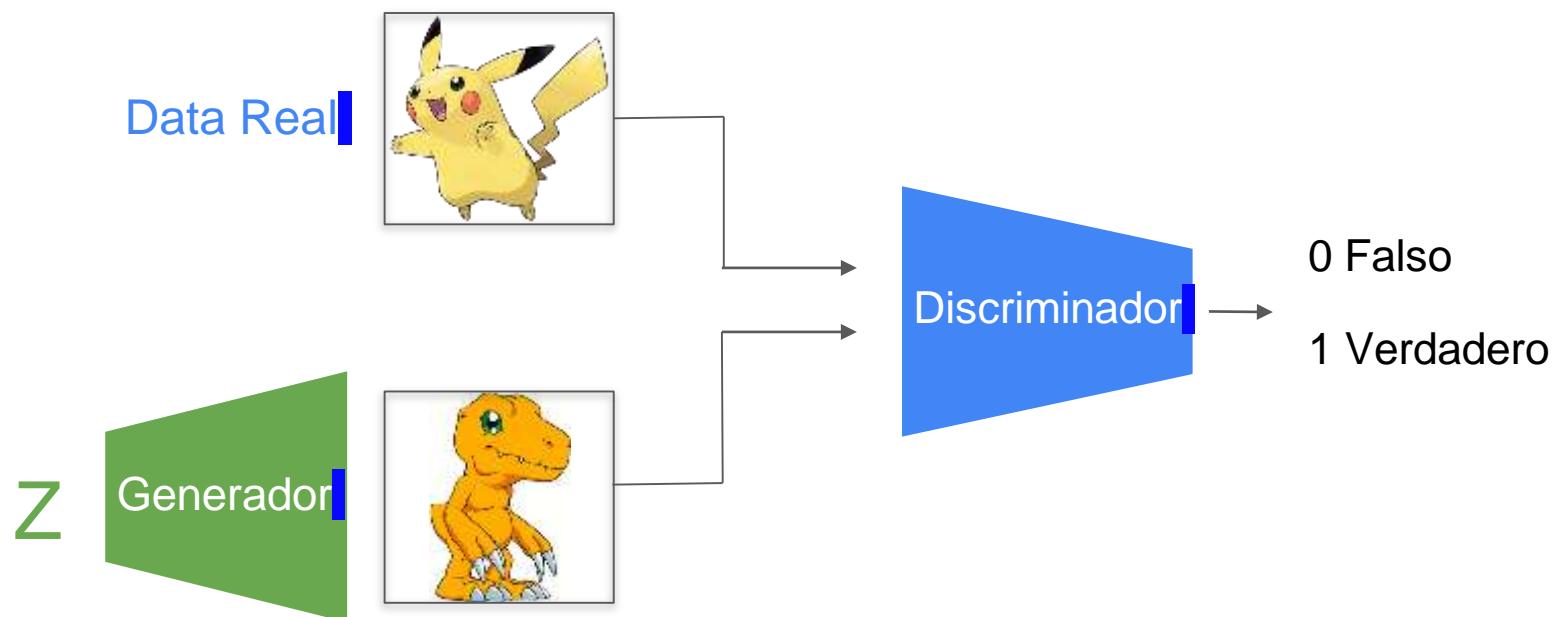


(c)

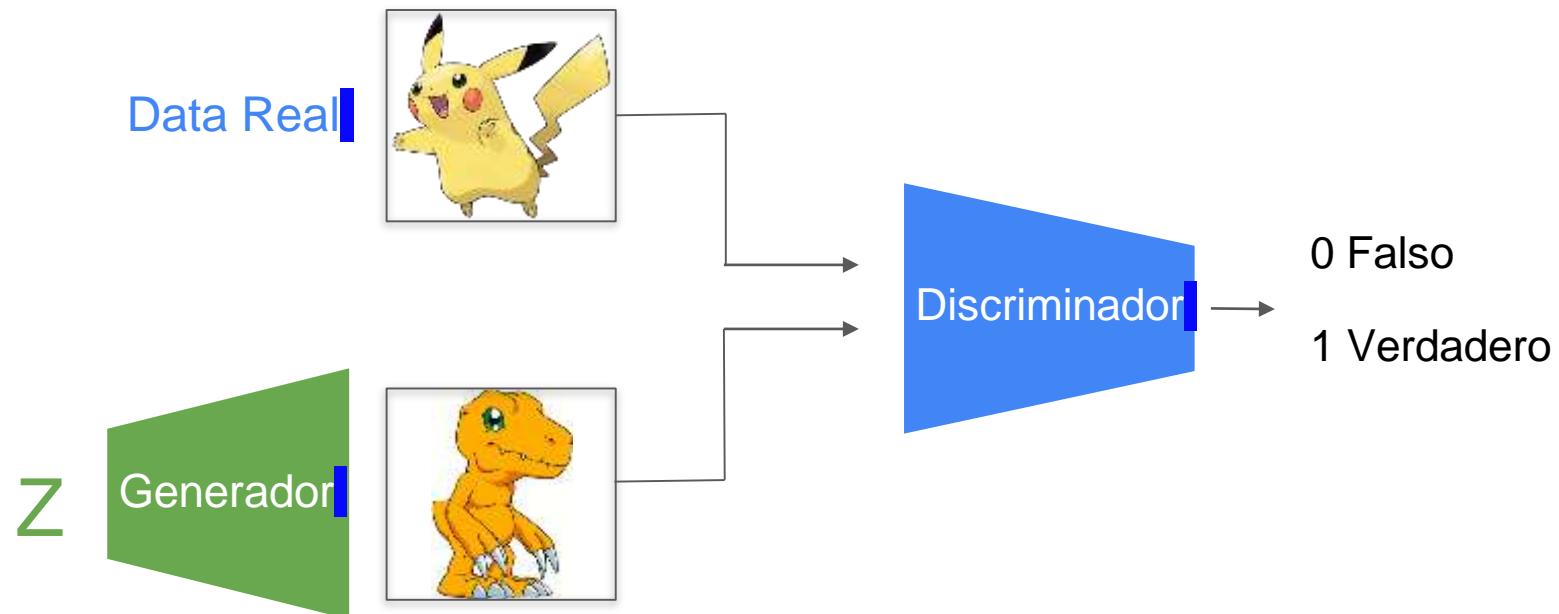


(d)

Función de pérdida

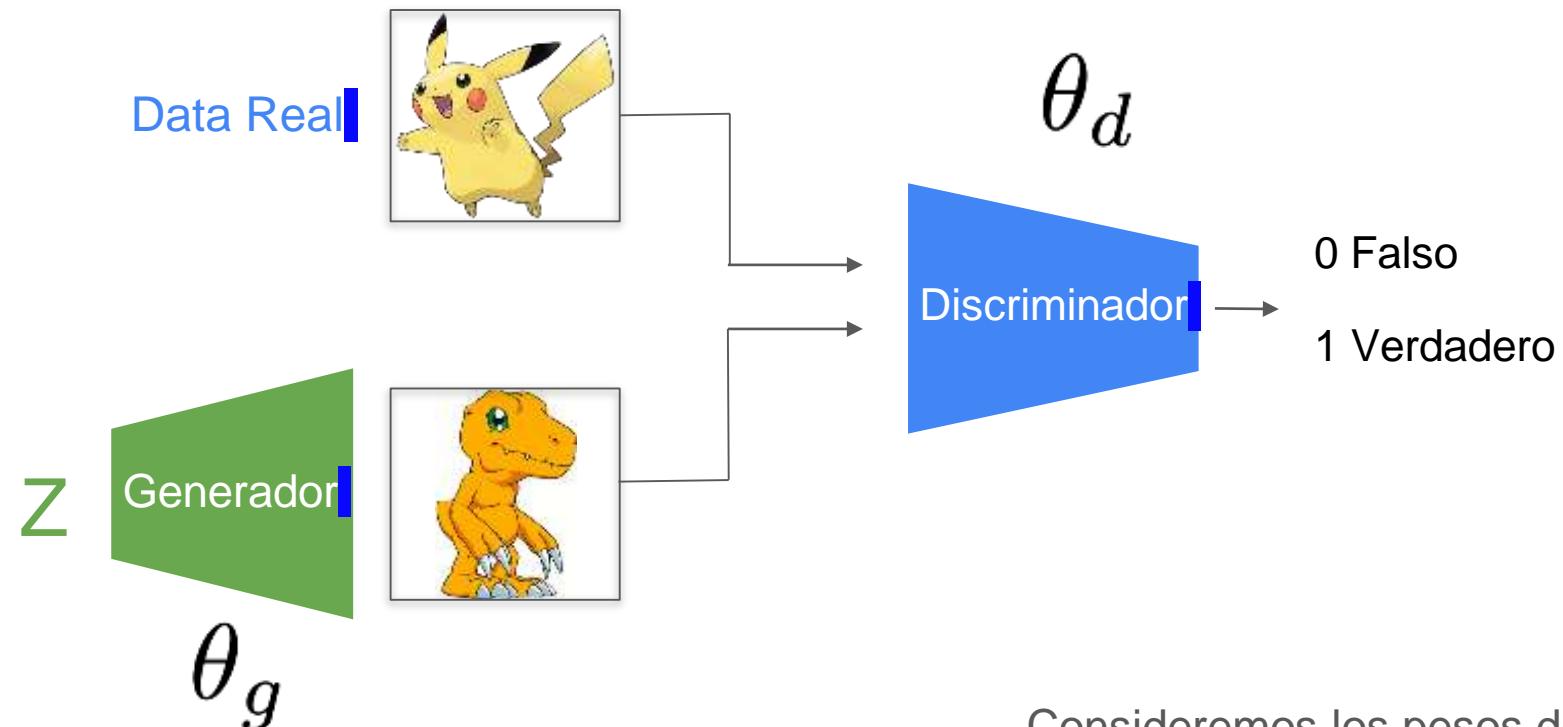


Función de pérdida



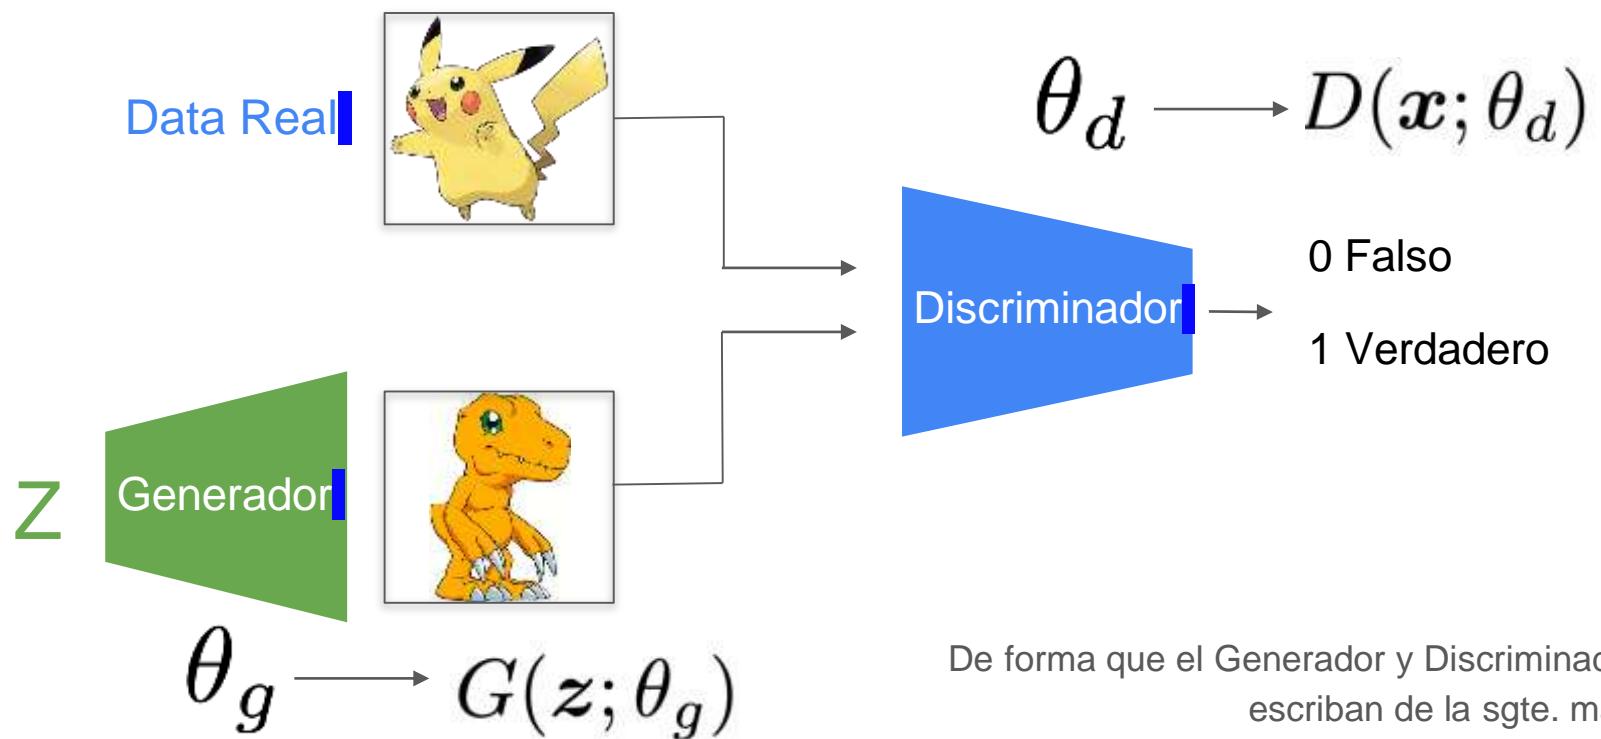
Consideremos los pesos de cada red de la siguiente forma...

Función de pérdida

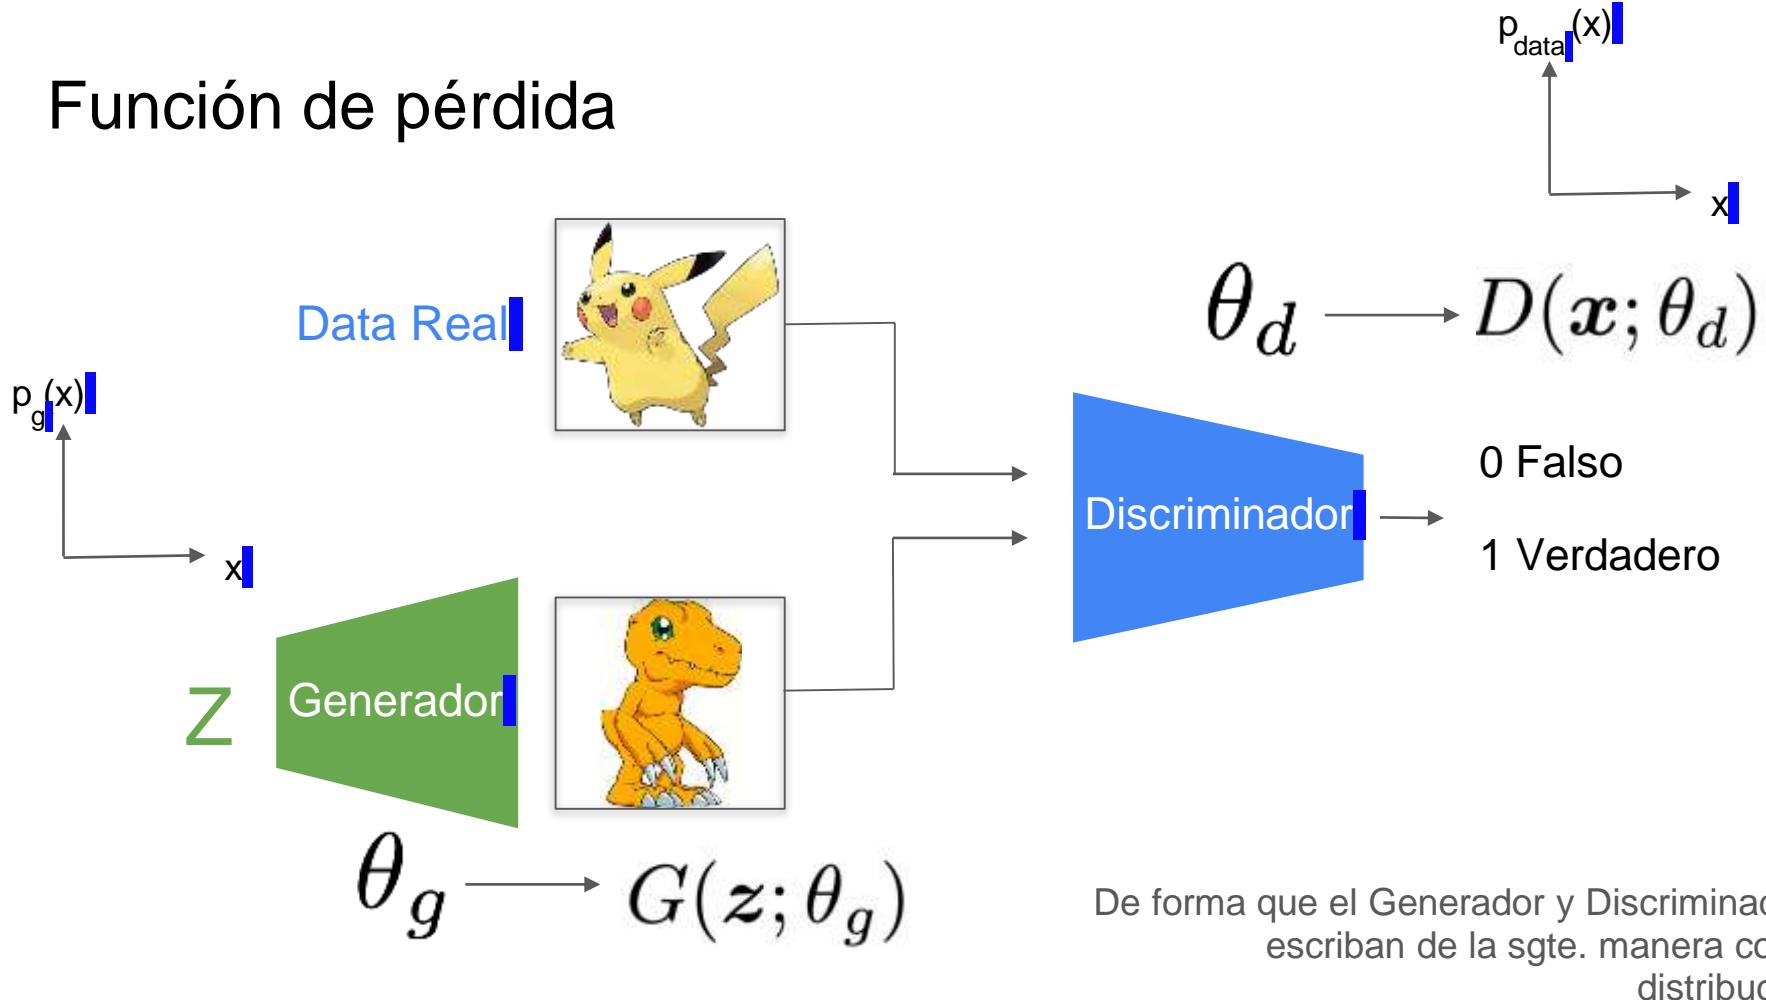


Consideremos los pesos de cada red de la siguiente forma...

Función de pérdida



Función de pérdida



De forma que el Generador y Discriminador se escriban de la sgte. manera con sus distribuciones

Función de pérdida

Función de valor



$$\min_G \max_D V(D, G) = \mathbb{E}_{\mathbf{x} \sim p_{\text{data}}(\mathbf{x})} [\log D(\mathbf{x})] + \mathbb{E}_{\mathbf{z} \sim p_{\mathbf{z}}(\mathbf{z})} [\log(1 - D(G(\mathbf{z})))].$$

De manera que G y D juegan un
MinMax de 2 jugadores

Función de pérdida

$$-\frac{1}{N} \sum_{i=1}^N y_i \cdot \log(p(y_i)) + (1 - y_i) \cdot \log(1 - p(y_i))$$



Función de BCE

Esto es similar a la función de
Binary Cross Entropy

Función de pérdida

$$-\frac{1}{N} \sum_{i=1}^N y_i \cdot \log(p(y_i)) + (1-y_i) \cdot \log(1-p(y_i))$$



Función de BCE

Cuando $y=1$

$$\hat{y} = D(x) \Rightarrow L = \ln[Dx]$$

Cuando $y=0$

$$\hat{y} = D(G(z)) \Rightarrow L = \ln[1 - D/G(z)]$$

Asimismo, cuando analizamos los valores de $y=1$ o $y=0$, vemos que...

Función de pérdida

$$\mathcal{L} = \ln[D(x)] + \ln[1 - D(\zeta(z))]$$

De forma que la función de valor
queda de la sgte. forma

Función de pérdida

$$\mathcal{L} = \ln[D(x)] + \ln[1 - D(\zeta(z))]$$

Pero esto es para una única muestra.
Se debe considerar la Esperanza

Función de pérdida

$$E(x) = \sum x p(x) \rightarrow (1+2+3+4+5+6)$$



La Esperanza...

Función de pérdida

$$E(x) = \sum x p(x) \rightarrow \frac{1}{6} (1+2+3+4+5+6)$$



La Esperanza...

Función de pérdida

$$E(x) = \sum x p(x) \rightarrow 3.5$$



La Esperanza...

Función de pérdida

$$\sum p_{\text{data}}(x) \ln[D(x)] + \sum p_z(z) \ln[1 - D(G(z))]$$

Aplicando la Esperanza a nuestra función de valor

Función de pérdida

$$\int p_{\text{data}}(x) \ln[D(x)] dx + \int p_z(z) \ln[1 - D(G(z))] dz$$

Aplicando la Esperanza a nuestra función de valor

Función de pérdida

$$V(G, D) = E_{x \sim P_{\text{data}}} [\ln(D(x))] + E_{z \sim P_z} [\ln(1 - D(G(z)))]$$

Aplicando la Esperanza a nuestra función de valor

Función de pérdida

MINIMIZAR
 D



$$\frac{\partial}{\partial \theta_j} \frac{1}{m} \cdot [\ln [1 - D(G(z))]]$$

MAXIMIZAR
 G



$$\frac{\partial}{\partial \theta_g} \frac{1}{m} \left[\ln [1 - D(G(z))] \right]$$

Cant. de
muestras

El algoritmo de entrenamiento
quedaría de la siguiente forma

Función de pérdida

Algorithm 1 Minibatch stochastic gradient descent training of generative adversarial nets. The number of steps to apply to the discriminator, k , is a hyperparameter. We used $k = 1$, the least expensive option, in our experiments.

for number of training iterations **do**

for k steps **do**

- Sample minibatch of m noise samples $\{\mathbf{z}^{(1)}, \dots, \mathbf{z}^{(m)}\}$ from noise prior $p_g(\mathbf{z})$.
- Sample minibatch of m examples $\{\mathbf{x}^{(1)}, \dots, \mathbf{x}^{(m)}\}$ from data generating distribution $p_{\text{data}}(\mathbf{x})$.
- Update the discriminator by ascending its stochastic gradient:

$$\nabla_{\theta_d} \frac{1}{m} \sum_{i=1}^m \left[\log D(\mathbf{x}^{(i)}) + \log (1 - D(G(\mathbf{z}^{(i)}))) \right].$$

end for

- Sample minibatch of m noise samples $\{\mathbf{z}^{(1)}, \dots, \mathbf{z}^{(m)}\}$ from noise prior $p_g(\mathbf{z})$.
- Update the generator by descending its stochastic gradient:

$$\nabla_{\theta_g} \frac{1}{m} \sum_{i=1}^m \log (1 - D(G(\mathbf{z}^{(i)}))).$$

end for

The gradient-based updates can use any standard gradient-based learning rule. We used momentum in our experiments.

El algoritmo de entrenamiento
quedaría de la siguiente forma

Aplicaciones

Crear personajes de anime

<https://arxiv.org/pdf/1708.05509.pdf>



Figure 7: Generated samples

Aplicaciones

Crear personajes de anime

<https://arxiv.org/pdf/1708.05509.pdf>

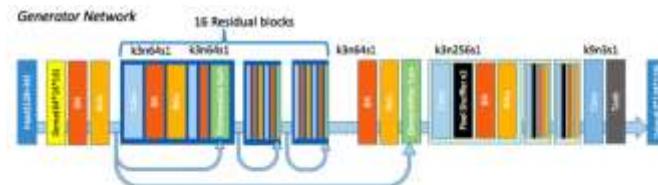


Figure 3: Generator Architecture



Figure 4: Discriminator Architecture

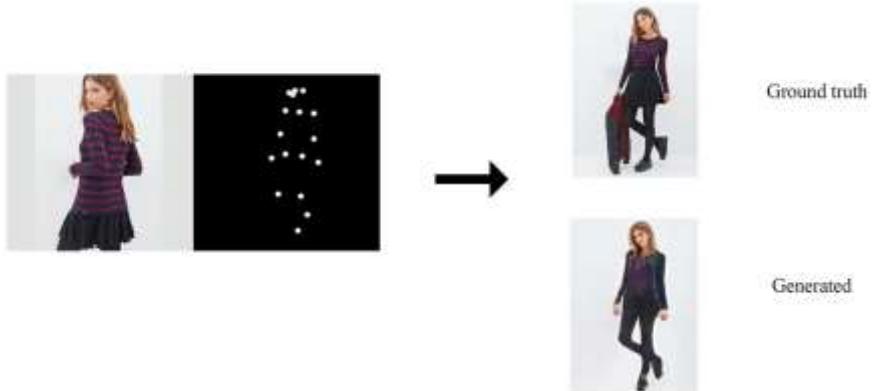


Figure 7: Generated samples

Aplicaciones

Personas posando

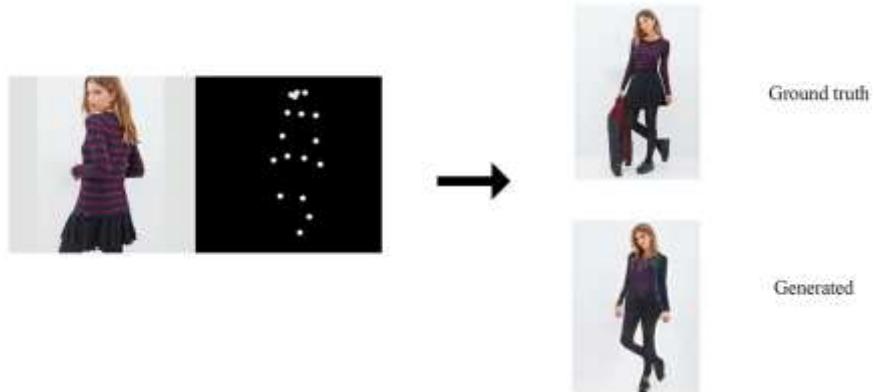
<https://arxiv.org/pdf/1705.09368.pdf>



Aplicaciones

Personas posando

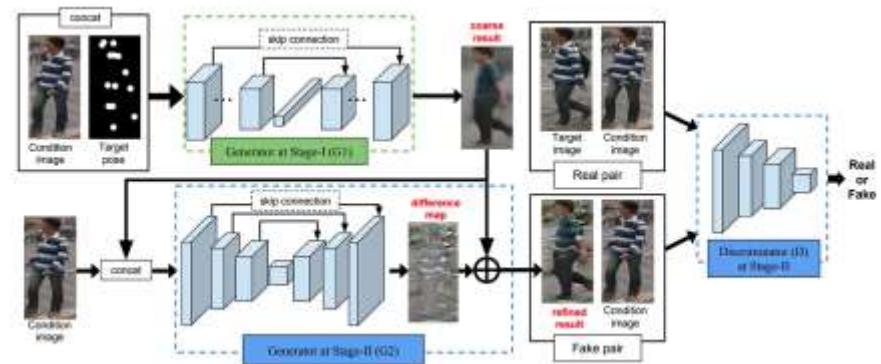
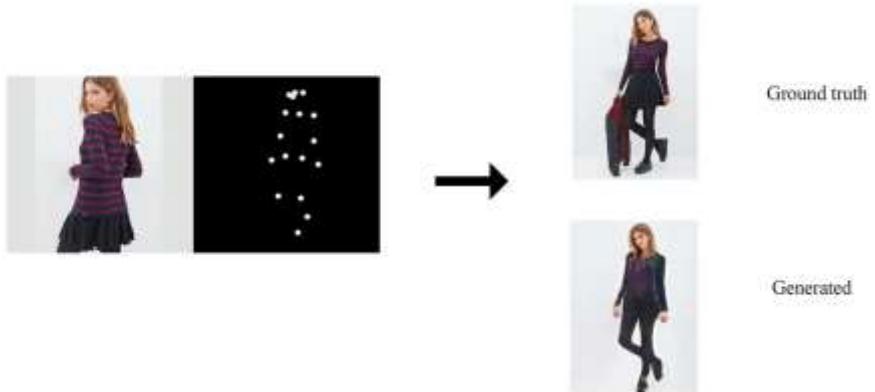
<https://arxiv.org/pdf/1705.09368.pdf>



Aplicaciones

Personas posando

<https://arxiv.org/pdf/1705.09368.pdf>



Aplicaciones

CycleGAN

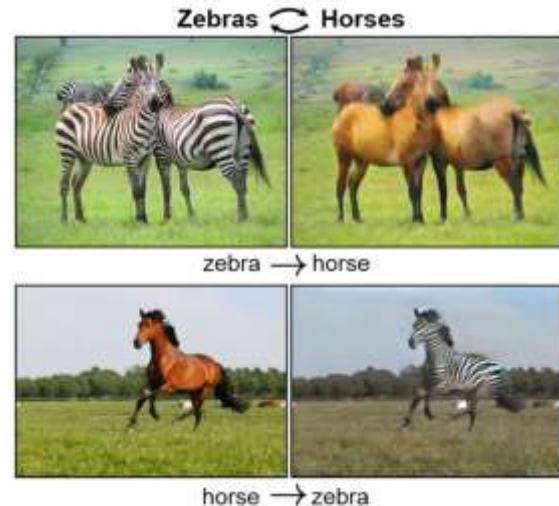
<https://github.com/junyanz/CycleGAN>



Aplicaciones

CycleGAN

<https://github.com/junyanz/CycleGAN>



Aplicaciones

StarGAN

<https://arxiv.org/pdf/1711.09020.pdf>

StarGAN es una traducción de imagen
a imagen de un dominio a otro.

Por ejemplo, dada una cara feliz,
queremos transformarla en una cara de
miedo.

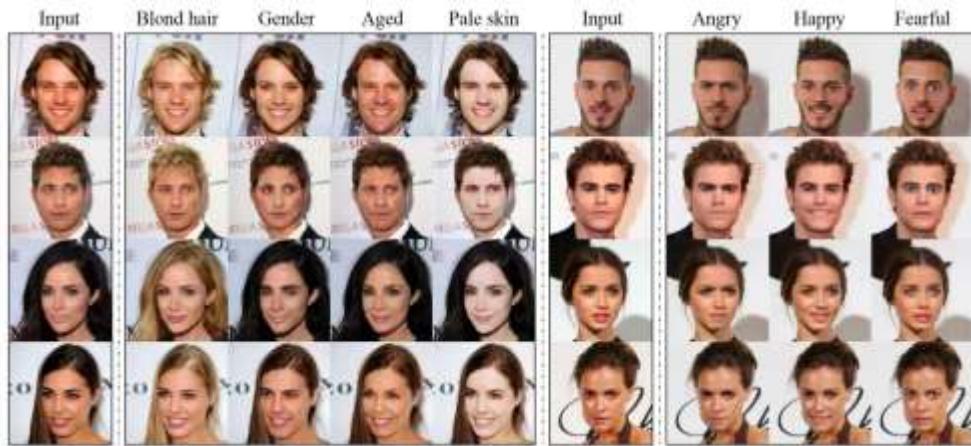
Aplicaciones

StarGAN

<https://arxiv.org/pdf/1711.09020.pdf>

StarGAN es una traducción de imagen a imagen de un dominio a otro.

Por ejemplo, dada una cara feliz, queremos transformarla en una cara de miedo.

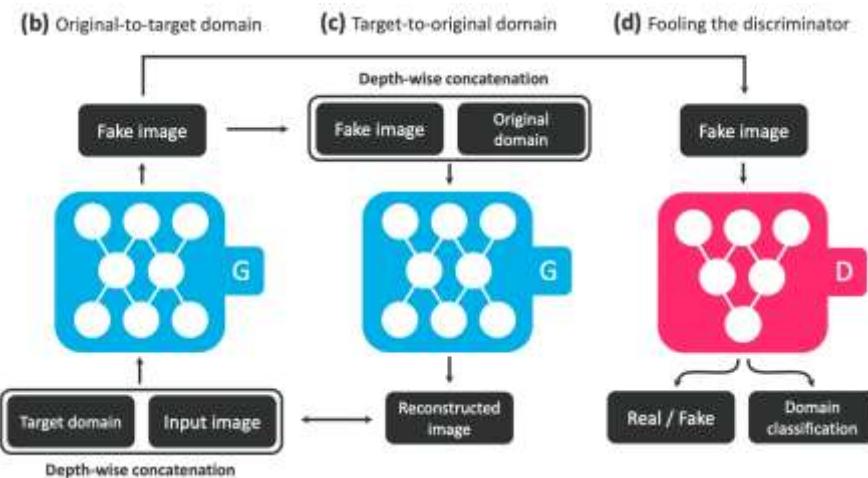


Aplicaciones

StarGAN

<https://arxiv.org/pdf/1711.09020.pdf>

En (b), el generador genera una imagen falsa basada en una imagen de entrada y una etiqueta de dominio de destino (por ejemplo, enojado).

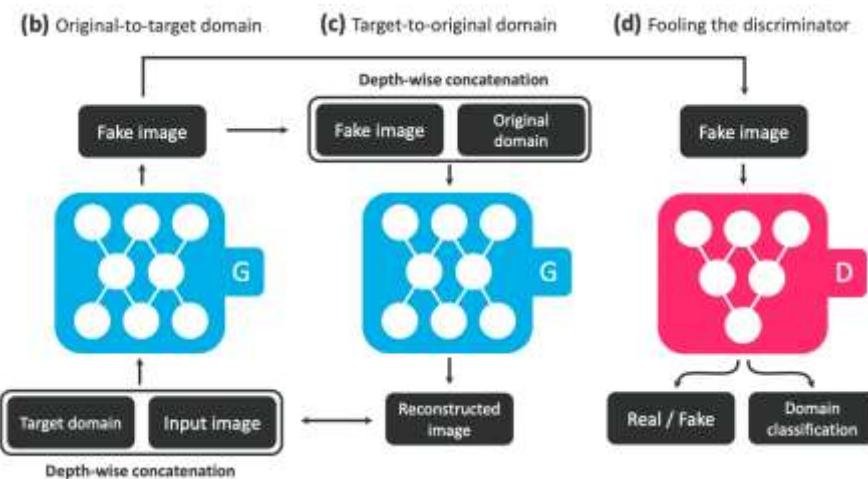


Aplicaciones

StarGAN

<https://arxiv.org/pdf/1711.09020.pdf>

En (d), alimentamos imágenes reales y falsas al discriminador para etiquetarlo como real o no, así como su clasificación de dominio.

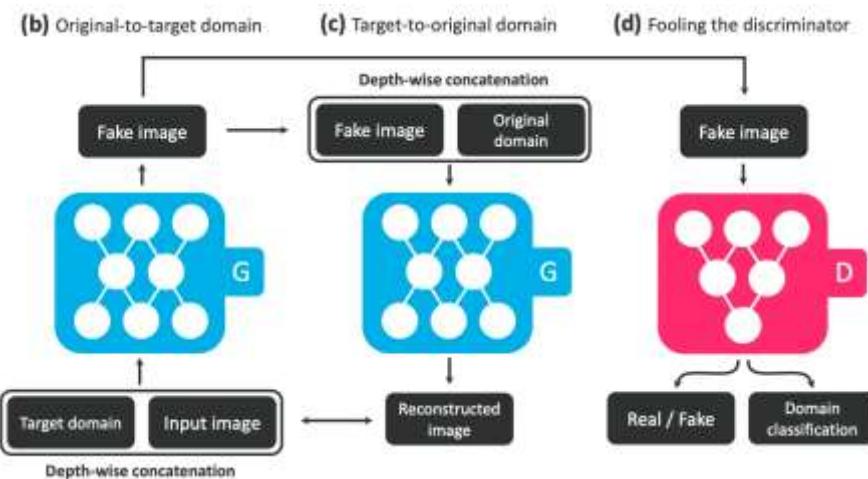


Aplicaciones

StarGAN

<https://arxiv.org/pdf/1711.09020.pdf>

La función de costo involucrará errores de reconstrucción así como el costo del discriminador en la identificación de las imágenes y sus etiquetas.



Aplicaciones

PixelDTGAN

<https://github.com/fxia22/PixelDTGAN>

Aplicaciones

PixelDTGAN

<https://github.com/fxia22/PixelDTGAN>



A source image.

Possible target images.

Aplicaciones

PixelDTGAN

<https://github.com/fxia22/PixelDTGAN>

Example results on LOOKBOOK dataset (top), left is input, right is generated clothes. Results on a similar dataset (bottom).
More results will be added soon.



Aplicaciones

PixelDTGAN

<https://github.com/fxia22/PixelDTGAN>

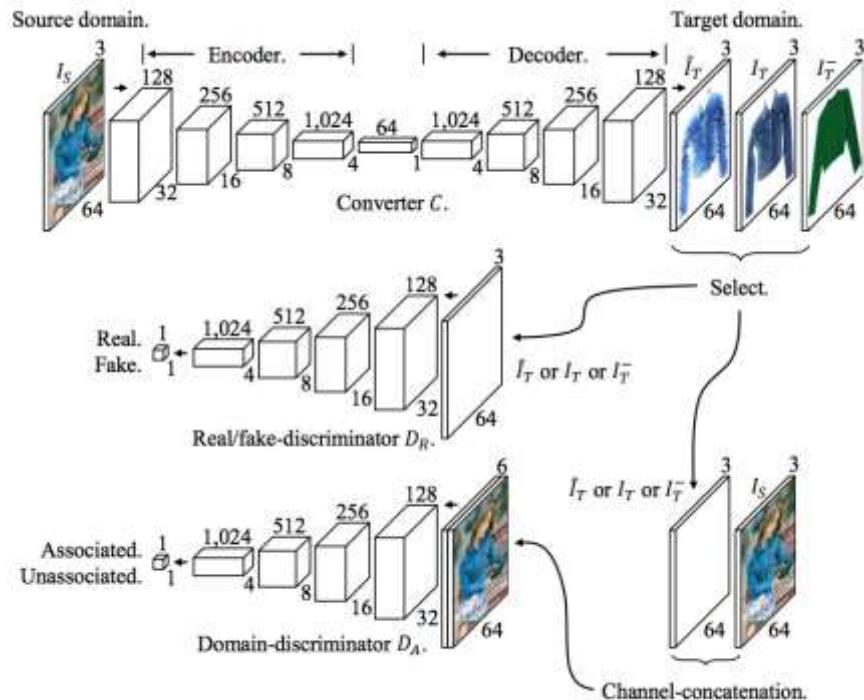


Fig. 2. Whole architecture for pixel-level domain transfer.

Aplicaciones

SuperResolution GAN

<https://arxiv.org/pdf/1609.04802.pdf>



Figure 2: From left to right: bicubic interpolation, deep residual network optimized for MSE, deep residual generative adversarial network optimized for a loss more sensitive to human perception, original HR image. Corresponding PSNR and SSIM are shown in brackets. [4× upscaling]

Aplicaciones

StyleGAN 2

<https://arxiv.org/pdf/1912.04958.pdf>



Aplicaciones

Text 2 Imagen GAN (StackGAN)

<https://github.com/hanzhanggit/StackGAN>

This flower has long thin yellow petals and a lot of yellow anthers in the center



Aplicaciones

Text 2 Imagen GAN (StackGAN)

<https://github.com/hanzhanggit/StackGAN>

A white bird with a black crown and yellow beak

Stage-I

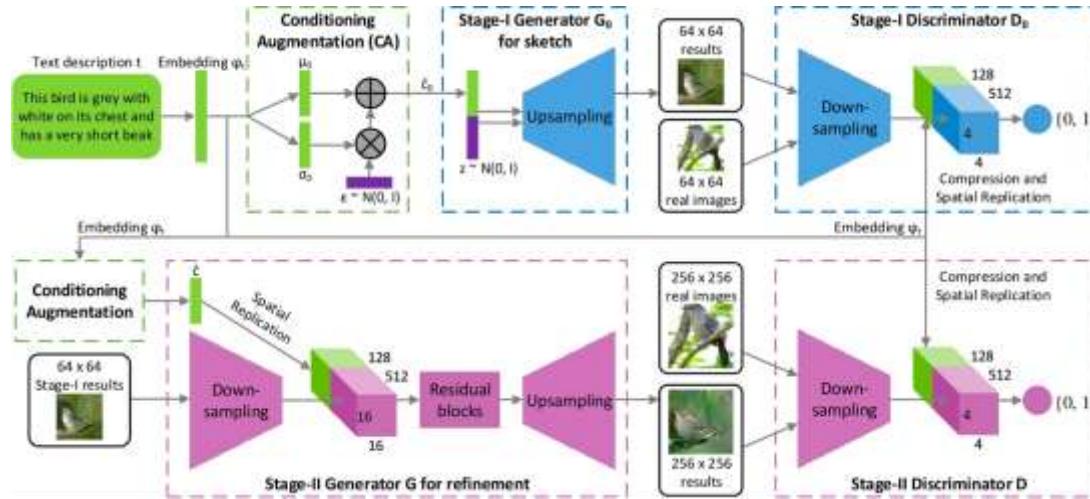


Stage-II

Aplicaciones

Text 2 Imagen GAN (StackGAN)

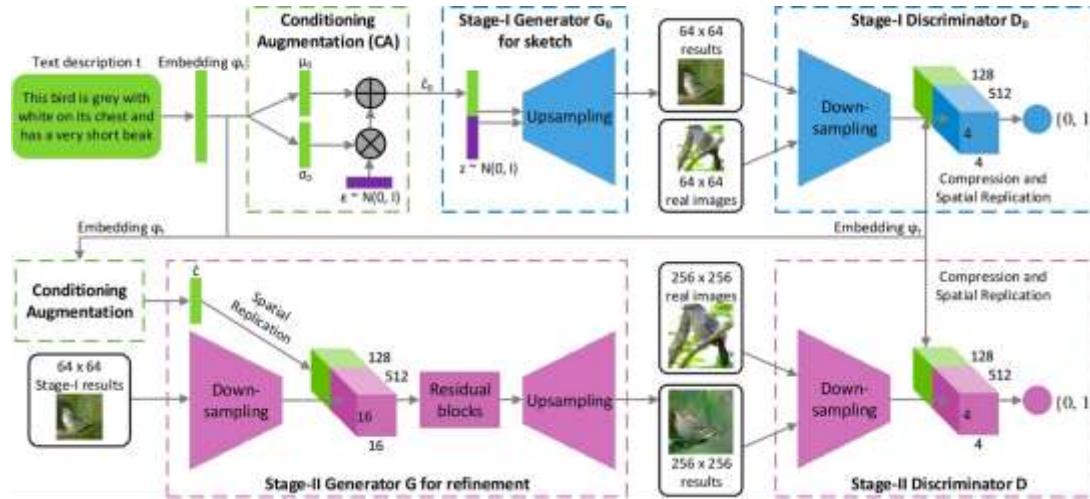
<https://github.com/hanzhanggit/StackGAN>



Aplicaciones

Text 2 Imagen GAN (StackGAN)

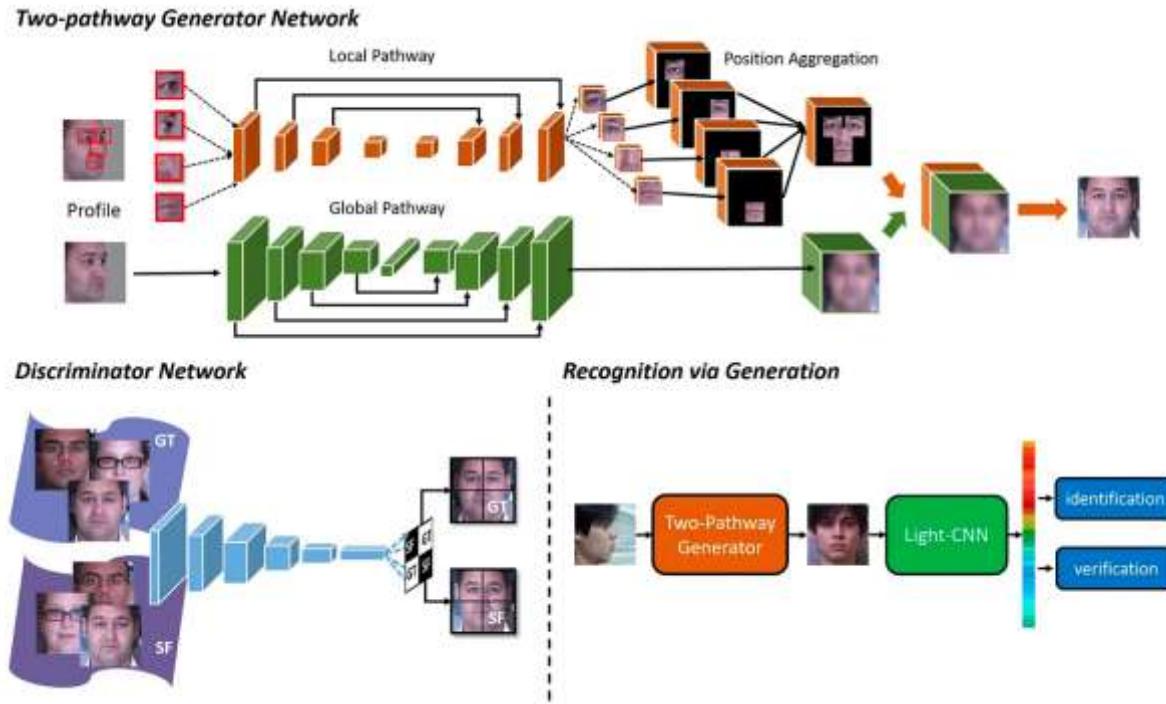
<https://github.com/hanzhanggit/StackGAN> - Demo: <https://www.craiyon.com/>



Aplicaciones

TPGAN

<https://arxiv.org/pdf/1704.04086.pdf>

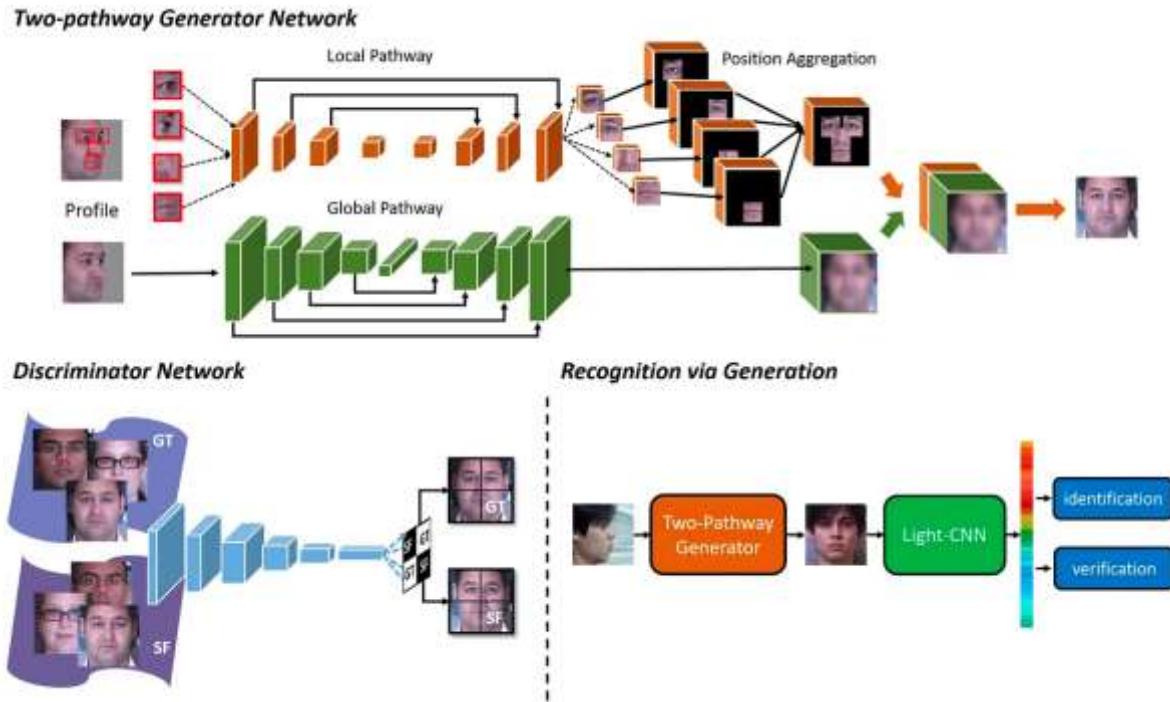


Aplicaciones

TPGAN

<https://arxiv.org/pdf/1704.04086.pdf>

Caras de síntesis en diferentes poses: Con una única imagen de entrada, crea caras en diferentes ángulos de visión.

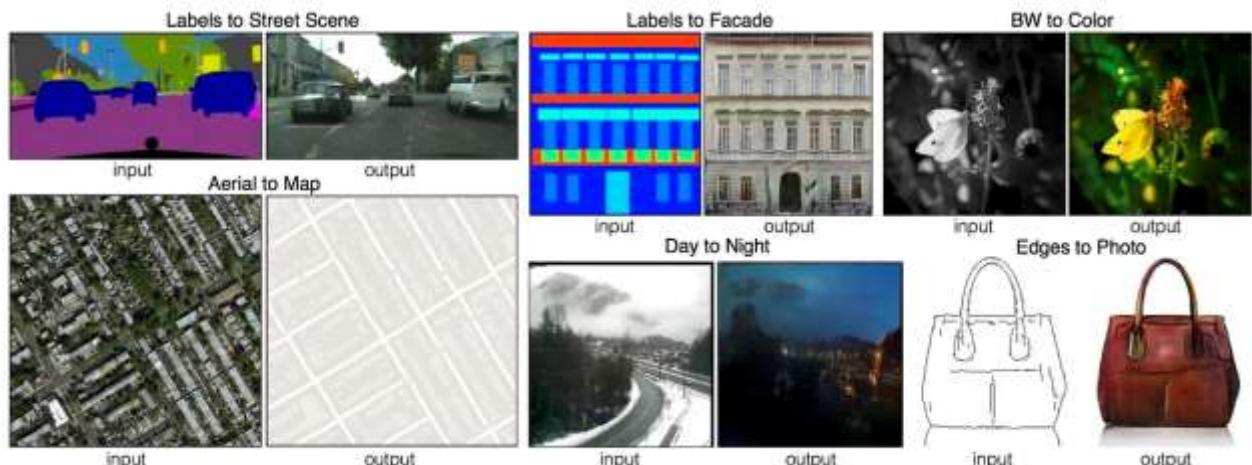


Aplicaciones

Pix2Pix

<https://github.com/phillipi/pix2pix>

Traducción de imágenes en diferentes tareas.



Aplicaciones

DeBlur

<https://arxiv.org/pdf/1711.07064.pdf>



Figure 2: GoPro images [25] processed by DeblurGAN. Blurred – left, DeblurGAN – center, ground truth sharp – right.

Aplicaciones

Neural Photo Editor

[https://github.com/ajbrock/
Neural-Photo-Editor](https://github.com/ajbrock/Neural-Photo-Editor)



Aplicaciones

MediGAN

<https://arxiv.org/pdf/1703.05921.pdf>

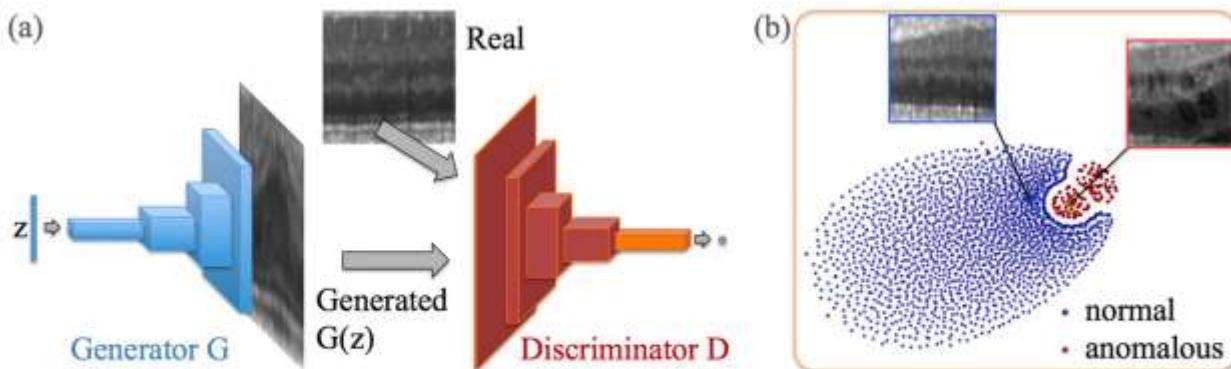
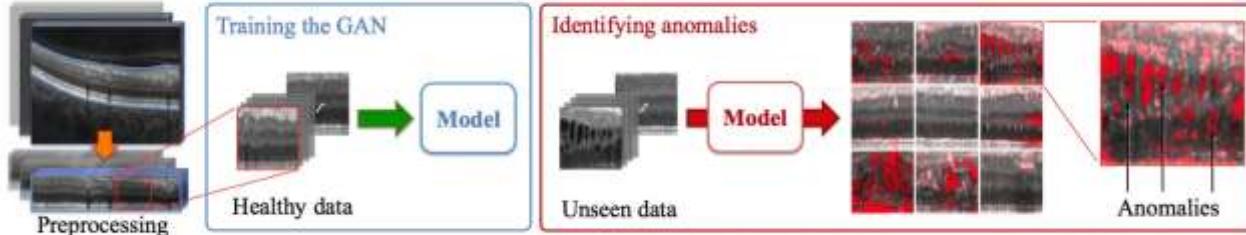


Fig. 2. (a) Deep convolutional generative adversarial network. (b) t-SNE embedding of normal (blue) and anomalous (red) images on the feature representation of the last convolution layer (orange in (a)) of the discriminator.

Aplicaciones

3D GAN

https://proceedings.neurips.cc/paper_files/paper/2016/file/44f683a84163b3523afe57c2e008bc8c-Paper.pdf



Aplicaciones

3D GAN

https://proceedings.neurips.cc/paper_files/paper/2016/file/44f683a84163b3523afe57c2e008bc8c-Paper.pdf

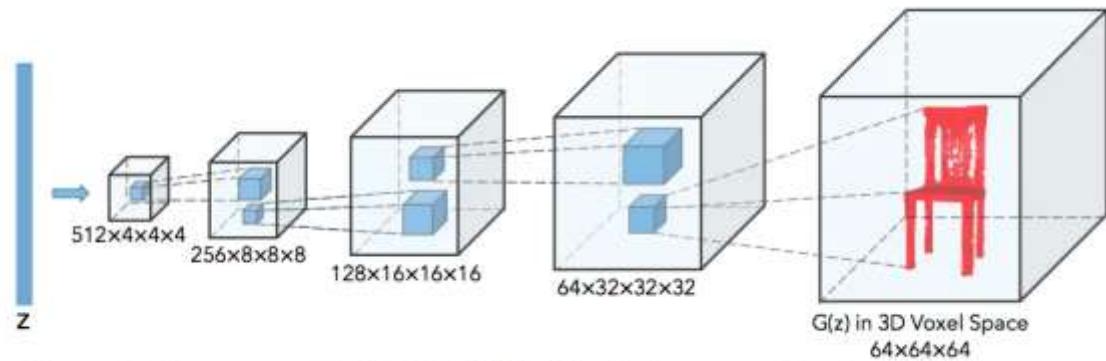


Figure 1: The generator in 3D-GAN. The discriminator mostly mirrors the generator.

Fuentes:

GAN (2014) Ian J. Goodfellow, Jean Pouget-Abadie, Mehdi Mirza, Bing Xu, David Warde-Farley, Sherjil Ozair, Aaron Courville, Yoshua Bengio

<https://arxiv.org/abs/1406.2661>

<https://github.com/soumith/ganhacks>

<https://jonathan-hui.medium.com/gan-some-cool-applications-of-gans-4c9ecca35900>