

# Deep Learning en Imágenes con GANs y Modelos de Difusión - Parte III

Prof. Peter Montalvo

# Contenido

- SR-GAN
- Función de Pérdida
- Función de Pérdida del Discriminador
- Pixel Shuffle
- Parametrized RELu

# SRGAN

- SR-GAN

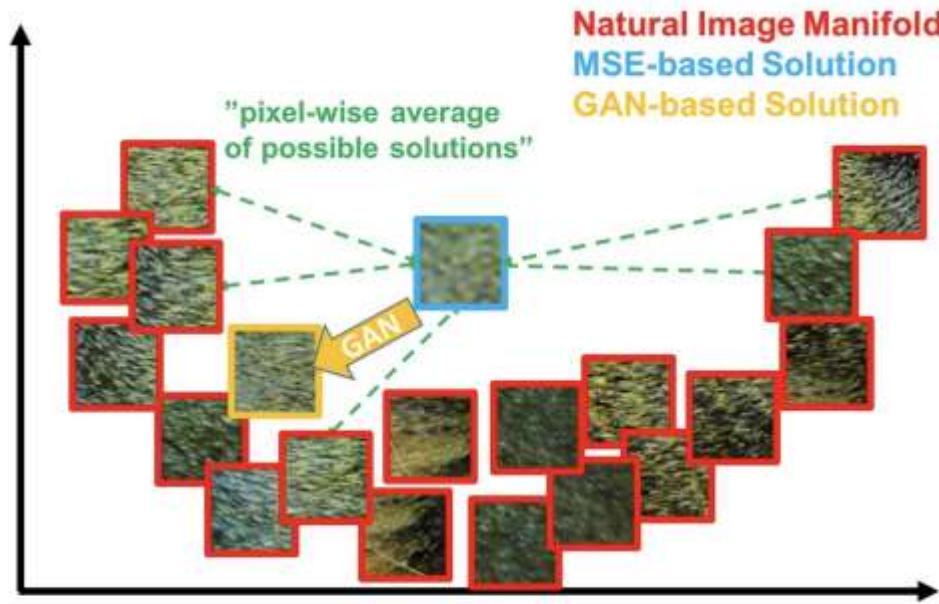
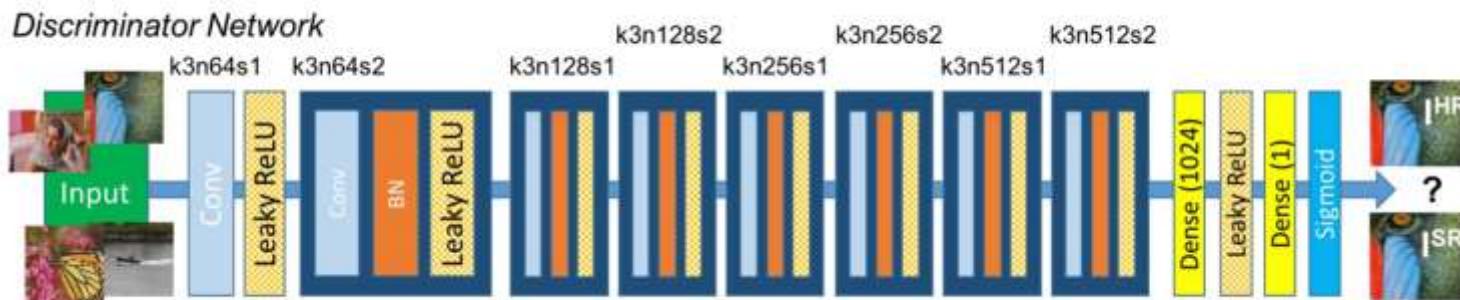
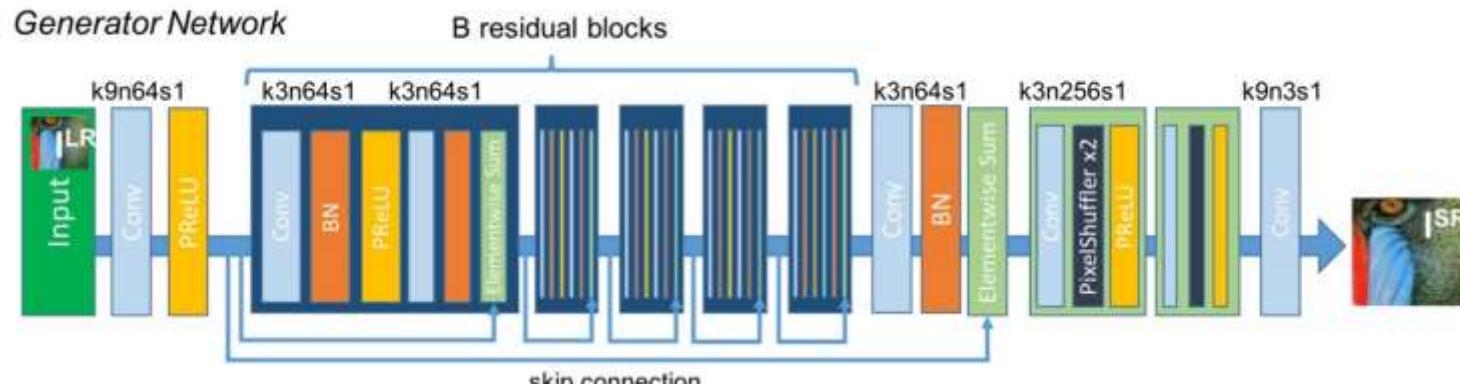


Figure 3: Illustration of patches from the natural image manifold (red) and super-resolved patches obtained with MSE (blue) and GAN (orange). The MSE-based solution appears overly smooth due to the pixel-wise average of possible solutions in the pixel space, while GAN drives the reconstruction towards the natural image manifold producing perceptually more convincing solutions.

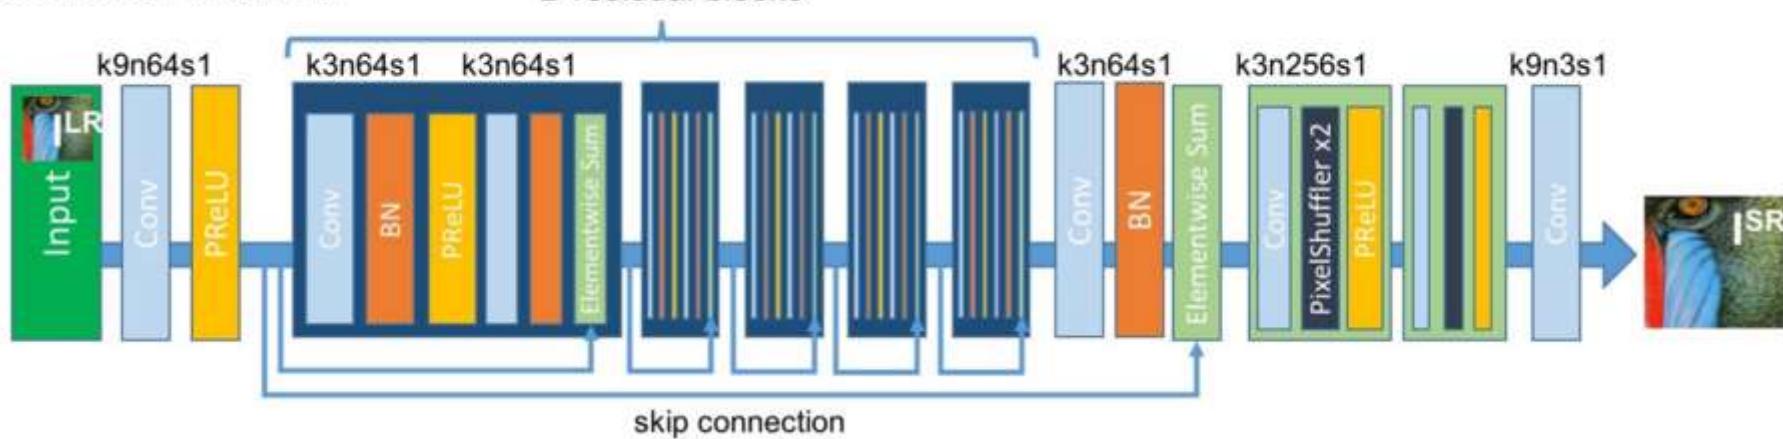
# SRGAN



# SRGAN

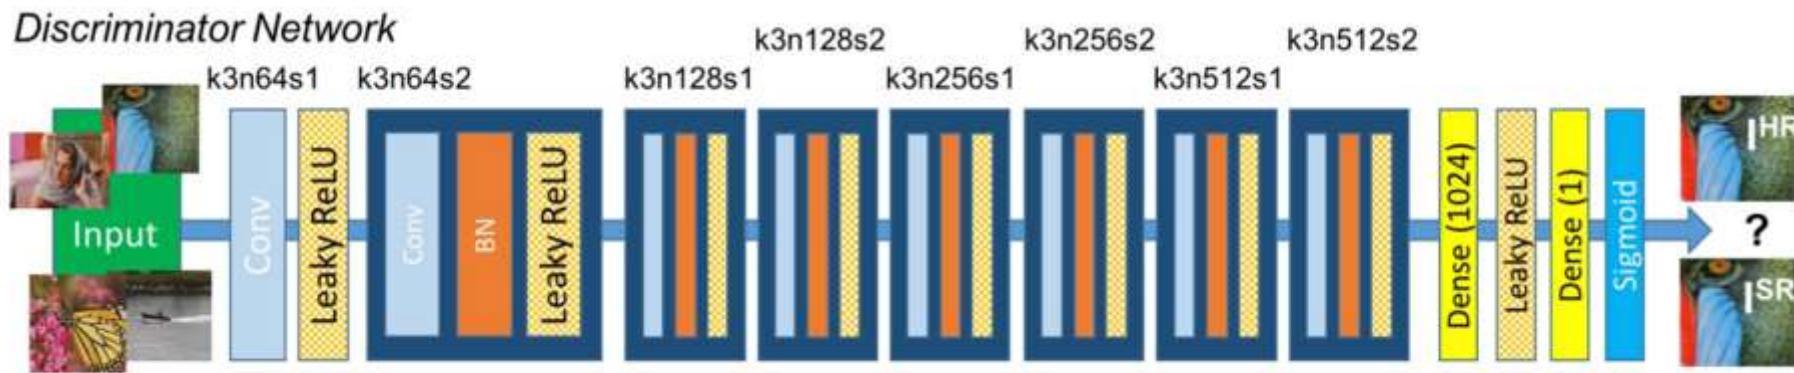
Generator Network

B residual blocks



# SRGAN

*Discriminator Network*



## SRGAN - Función de Pérdida

$$l^{SR} = \underbrace{l_X^{SR}}_{\substack{\text{content loss}}} + 10^{-3} \underbrace{l_{Gen}^{SR}}_{\substack{\text{adversarial loss}}} \\ \underbrace{\hspace{10em}}_{\text{perceptual loss (for VGG based content losses)}}$$

# SRGAN - Función de Pérdida

$$l_{MSE}^{SR} = \frac{1}{r^2WH} \sum_{x=1}^{rW} \sum_{y=1}^{rH} (I_{x,y}^{HR} - G_{\theta_G}(I^{LR})_{x,y})^2$$

Usualmente se calcula la diferencia entre la **Imagen en Alta Resolución** y el **resultado del Generador** con **la imagen en Baja Resolución**

# SRGAN - Función de Pérdida

$$l_{MSE}^{SR} = \frac{1}{r^2WH} \sum_{x=1}^{rW} \sum_{y=1}^{rH} (I_{x,y}^{HR} - G_{\theta_G}(I^{LR})_{x,y})^2$$

Este error también es parte del Peak signal-to-noise ratio...

Más info: [https://en.wikipedia.org/wiki/Peak\\_signal-to-noise\\_ratio](https://en.wikipedia.org/wiki/Peak_signal-to-noise_ratio)

# SRGAN - Función de Pérdida

$$MSE = \frac{1}{m n} \sum_{i=0}^{m-1} \sum_{j=0}^{n-1} [I(i, j) - K(i, j)]^2.$$

La relación señal-ruido máxima (PSNR) es un término de ingeniería para la relación entre la potencia máxima posible de una señal y la potencia del ruido corruptor que afecta la fidelidad de su representación.

$$\begin{aligned} PSNR &= 10 \cdot \log_{10} \left( \frac{MAX_I^2}{MSE} \right) \\ &= 20 \cdot \log_{10} \left( \frac{MAX_I}{\sqrt{MSE}} \right) \\ &= 20 \cdot \log_{10}(MAX_I) - 10 \cdot \log_{10}(MSE). \end{aligned}$$

Más info: [https://en.wikipedia.org/wiki/Peak\\_signal-to-noise\\_ratio](https://en.wikipedia.org/wiki/Peak_signal-to-noise_ratio)

# SRGAN - Función de Pérdida

$$MSE = \frac{1}{m n} \sum_{i=0}^{m-1} \sum_{j=0}^{n-1} [I(i, j) - K(i, j)]^2.$$

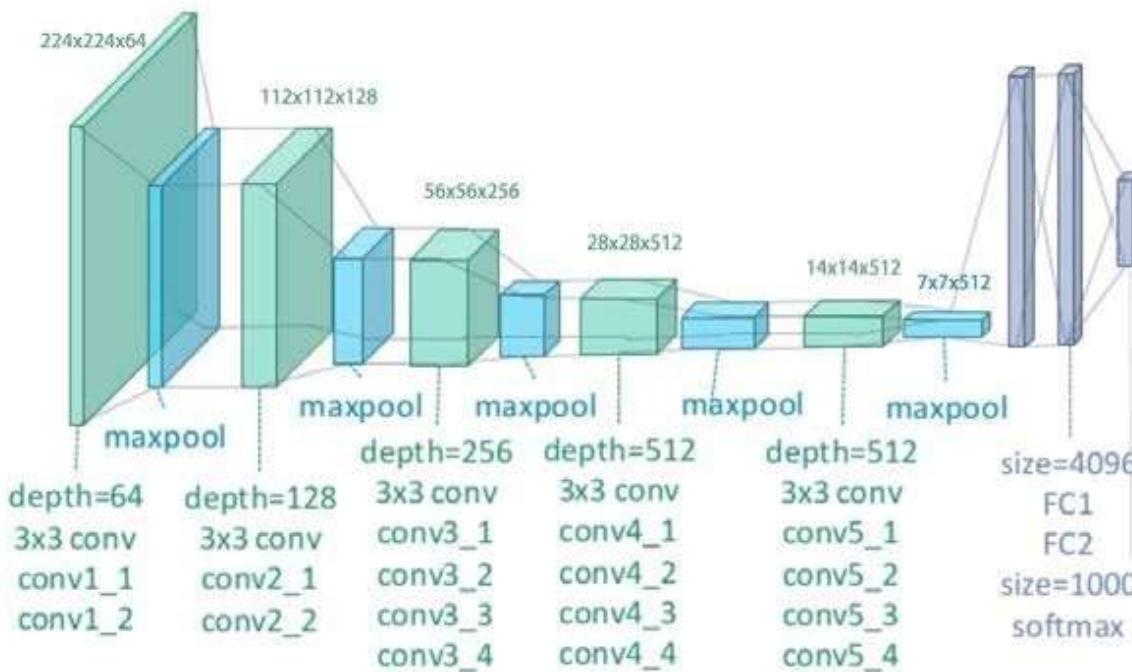
$$\begin{aligned} PSNR &= 10 \cdot \log_{10} \left( \frac{MAX_I^2}{MSE} \right) \\ &= 20 \cdot \log_{10} \left( \frac{MAX_I}{\sqrt{MSE}} \right) \\ &= 20 \cdot \log_{10}(MAX_I) - 10 \cdot \log_{10}(MSE). \end{aligned}$$

Debido a que muchas señales tienen un rango dinámico muy amplio, la PSNR generalmente se expresa como una cantidad logarítmica usando la escala de decibeles.

PSNR se usa comúnmente para cuantificar la calidad de reconstrucción de imágenes y videos sujetos a compresión con pérdidas.

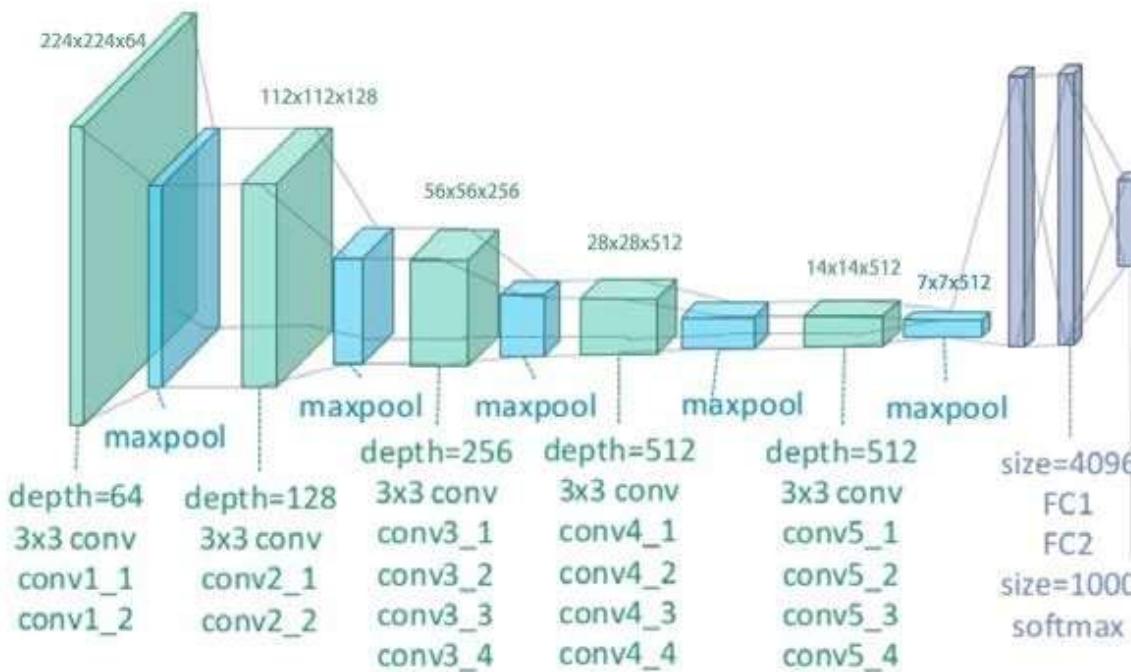
Más info: [https://en.wikipedia.org/wiki/Peak\\_signal-to-noise\\_ratio](https://en.wikipedia.org/wiki/Peak_signal-to-noise_ratio)

# SRGAN - Función de Pérdida



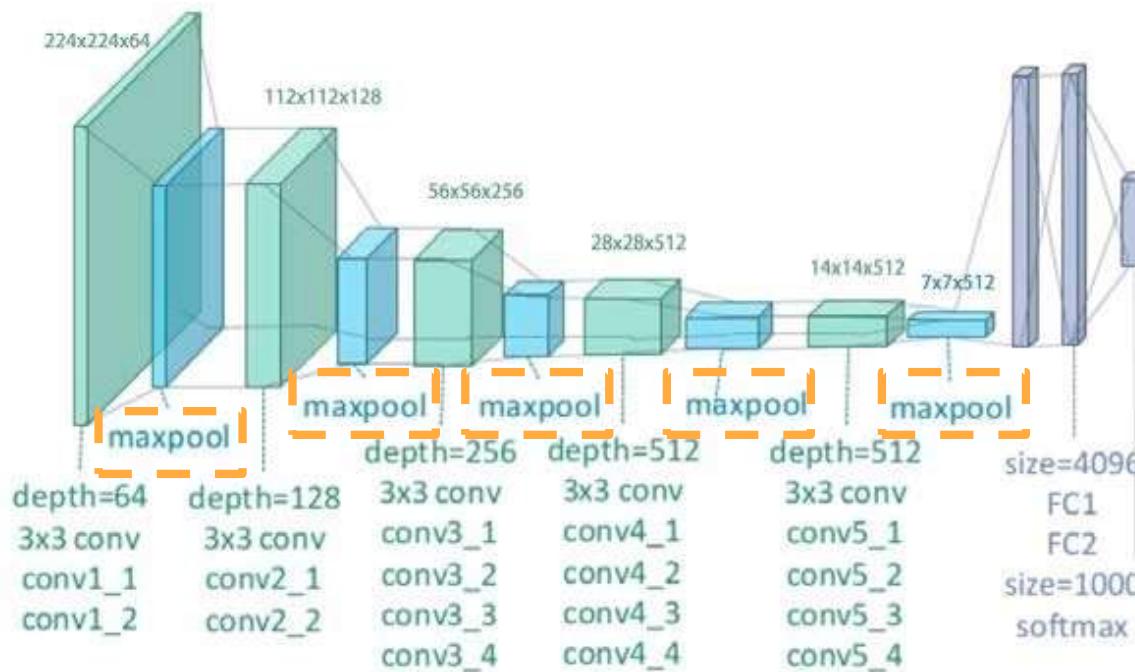
Consideremos que el discriminador es un VGG-19...

# SRGAN - Función de Pérdida



Consideremos que el discriminador es un VGG-19... cuya tarea es clasificar 1000 cat.

# SRGAN - Función de Pérdida



Consideremos cada maxpool como un **feature map** y será llamada  $\phi$

# SRGAN - Función de Pérdida

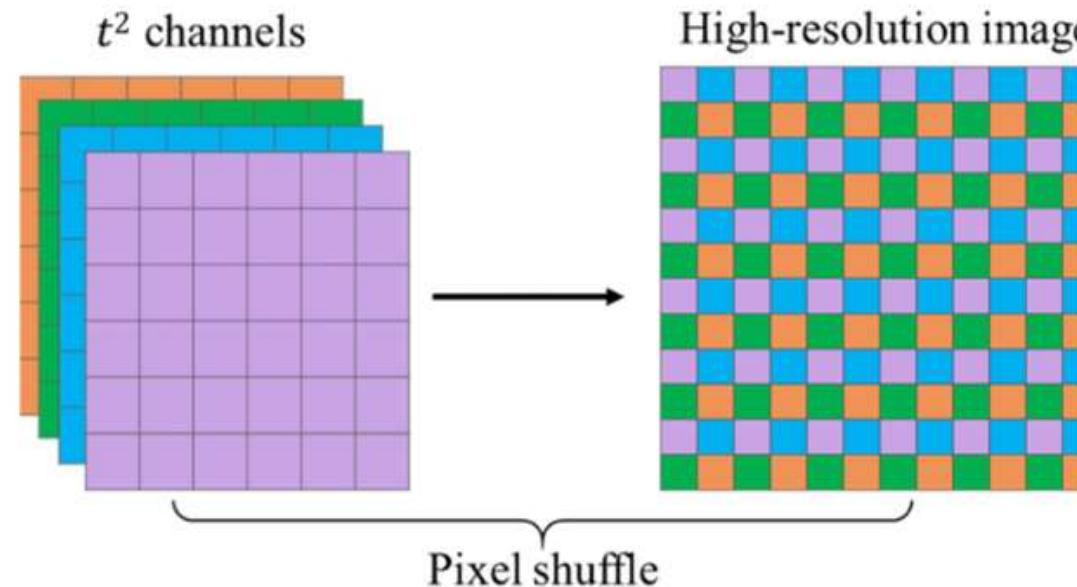
$$l_{VGG/i.j}^{SR} = \frac{1}{W_{i,j}H_{i,j}} \sum_{x=1}^{W_{i,j}} \sum_{y=1}^{H_{i,j}} (\phi_{i,j}(I^{HR})_{x,y} - \phi_{i,j}(G_{\theta_G}(I^{LR}))_{x,y})^2$$

De manera que el error es la diferencia del mapa de activación de la imagen en alta resolución vs. el mapa de activación de la imagen obtenida por el Generador con la imagen de Baja Resolución

# SRGAN - Función de Pérdida del Discriminador

$$l_{Gen}^{SR} = \sum_{n=1}^N -\log D_{\theta_D}(G_{\theta_G}(I^{LR}))$$

# Pixel Shuffle



Más info:

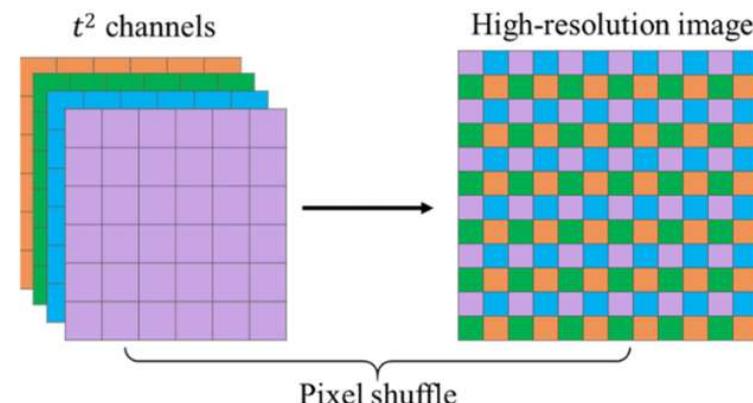
[https://paperswithcode.com/  
method/pixelshuffle](https://paperswithcode.com/method/pixelshuffle)

# Pixel Shuffle

**PixelShuffle** is an operation used in super-resolution models to implement efficient sub-pixel convolutions with a stride of  $1/r$ . Specifically it rearranges elements in a tensor of shape  $(*, C \times r^2, H, W)$  to a tensor of shape  $(*, C, H \times r, W \times r)$ .

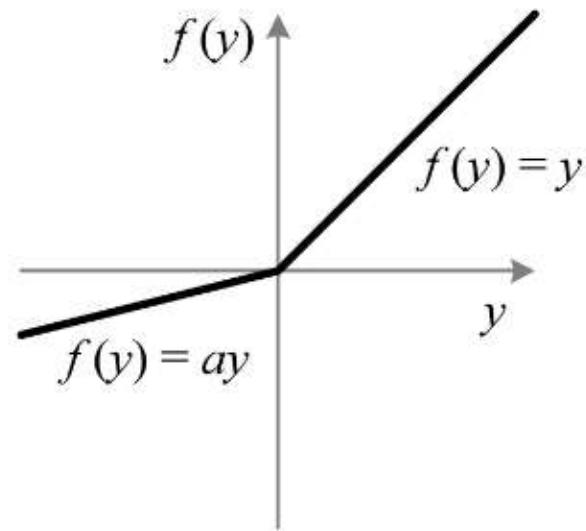
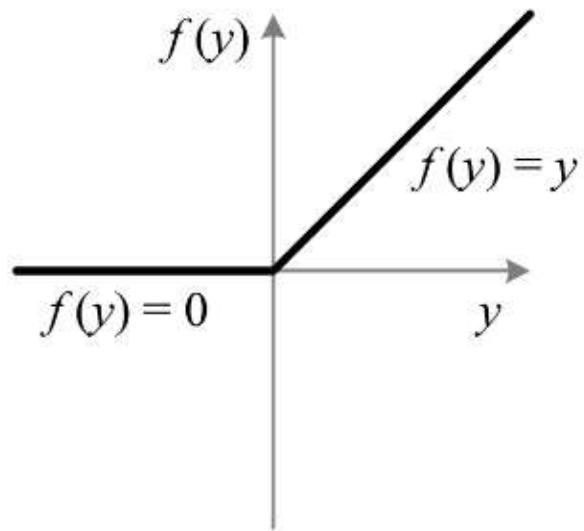
Image Source: [Remote Sensing Single-Image Resolution Improvement Using A Deep Gradient-Aware Network with Image-Specific Enhancement](#)

Source: [Real-Time Single Image and Video Super-Resolution Using an Efficient Sub-Pixel Convolutional Neural Network](#)



Más info:  
[https://paperswithcode.com/  
method/pixelshuffle](https://paperswithcode.com/method/pixelshuffle)

# Parametrized ReLu



Más info:

[https://paperswithcode.com/  
method/prelu](https://paperswithcode.com/method/prelu)

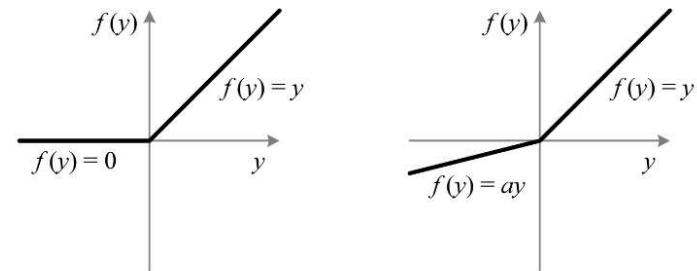
# Parametrized ReLU

A Parametric Rectified Linear Unit, or PReLU, is an activation function that generalizes the traditional rectified unit with a slope for negative values. Formally:

$$f(y_i) = y_i \text{ if } y_i \geq 0$$

$$f(y_i) = a_i y_i \text{ if } y_i \leq 0$$

The intuition is that different layers may require different types of nonlinearity. Indeed the authors find in experiments with convolutional neural networks that PReLUs for the initial layer have more positive slopes, i.e. closer to linear. Since the filters of the first layers are Gabor-like filters such as edge or texture detectors, this shows a circumstance where positive and negative responses of filters are respected. In contrast the authors find deeper layers have smaller coefficients, suggesting the model becomes more discriminative at later layers (while it wants to retain more information at earlier layers).



Más info:

[https://paperswithcode.com/  
method/prelu](https://paperswithcode.com/method/prelu)

# Fuentes

<https://arxiv.org/abs/1609.04802v5>